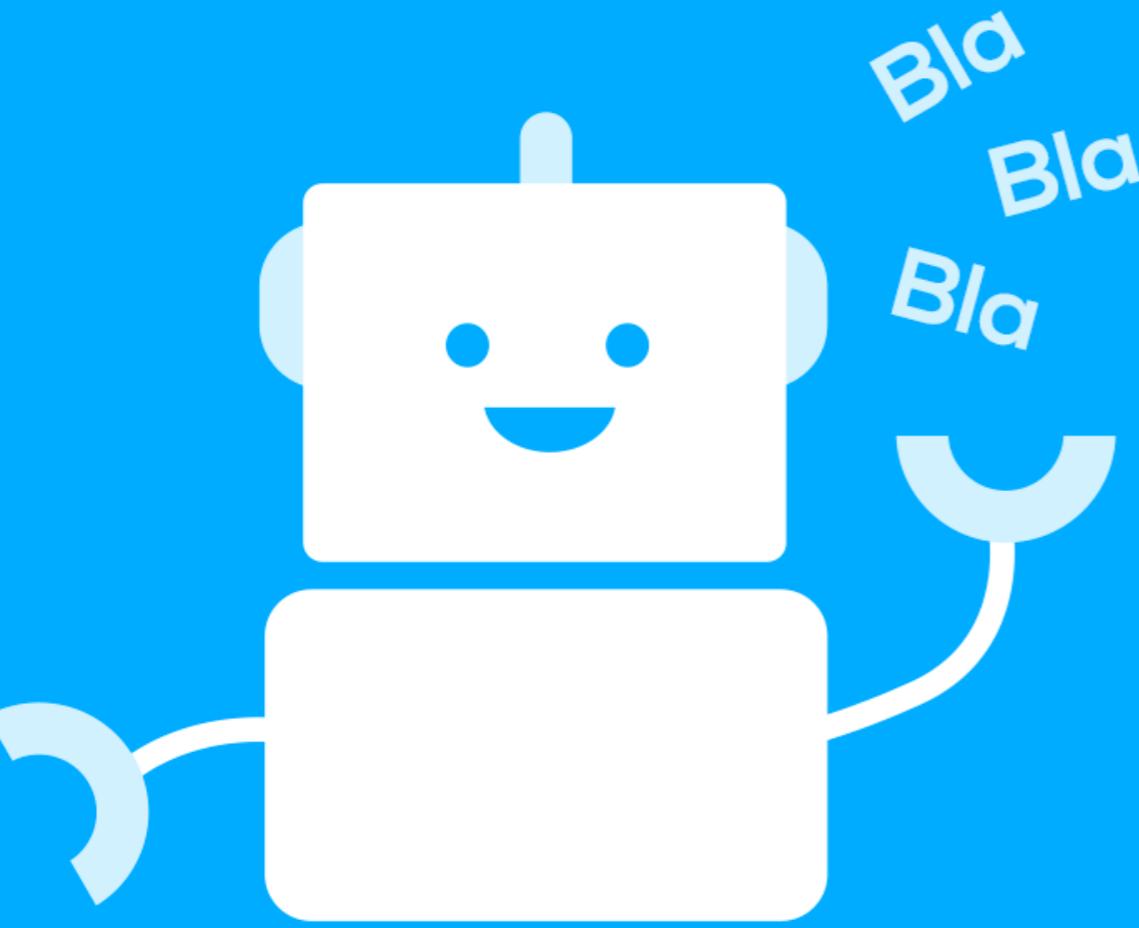


Building Task-Oriented Neural Dialog System

Head First Theory and Practice

Yanran Li
The Hong Kong Polytechnic University



vs.

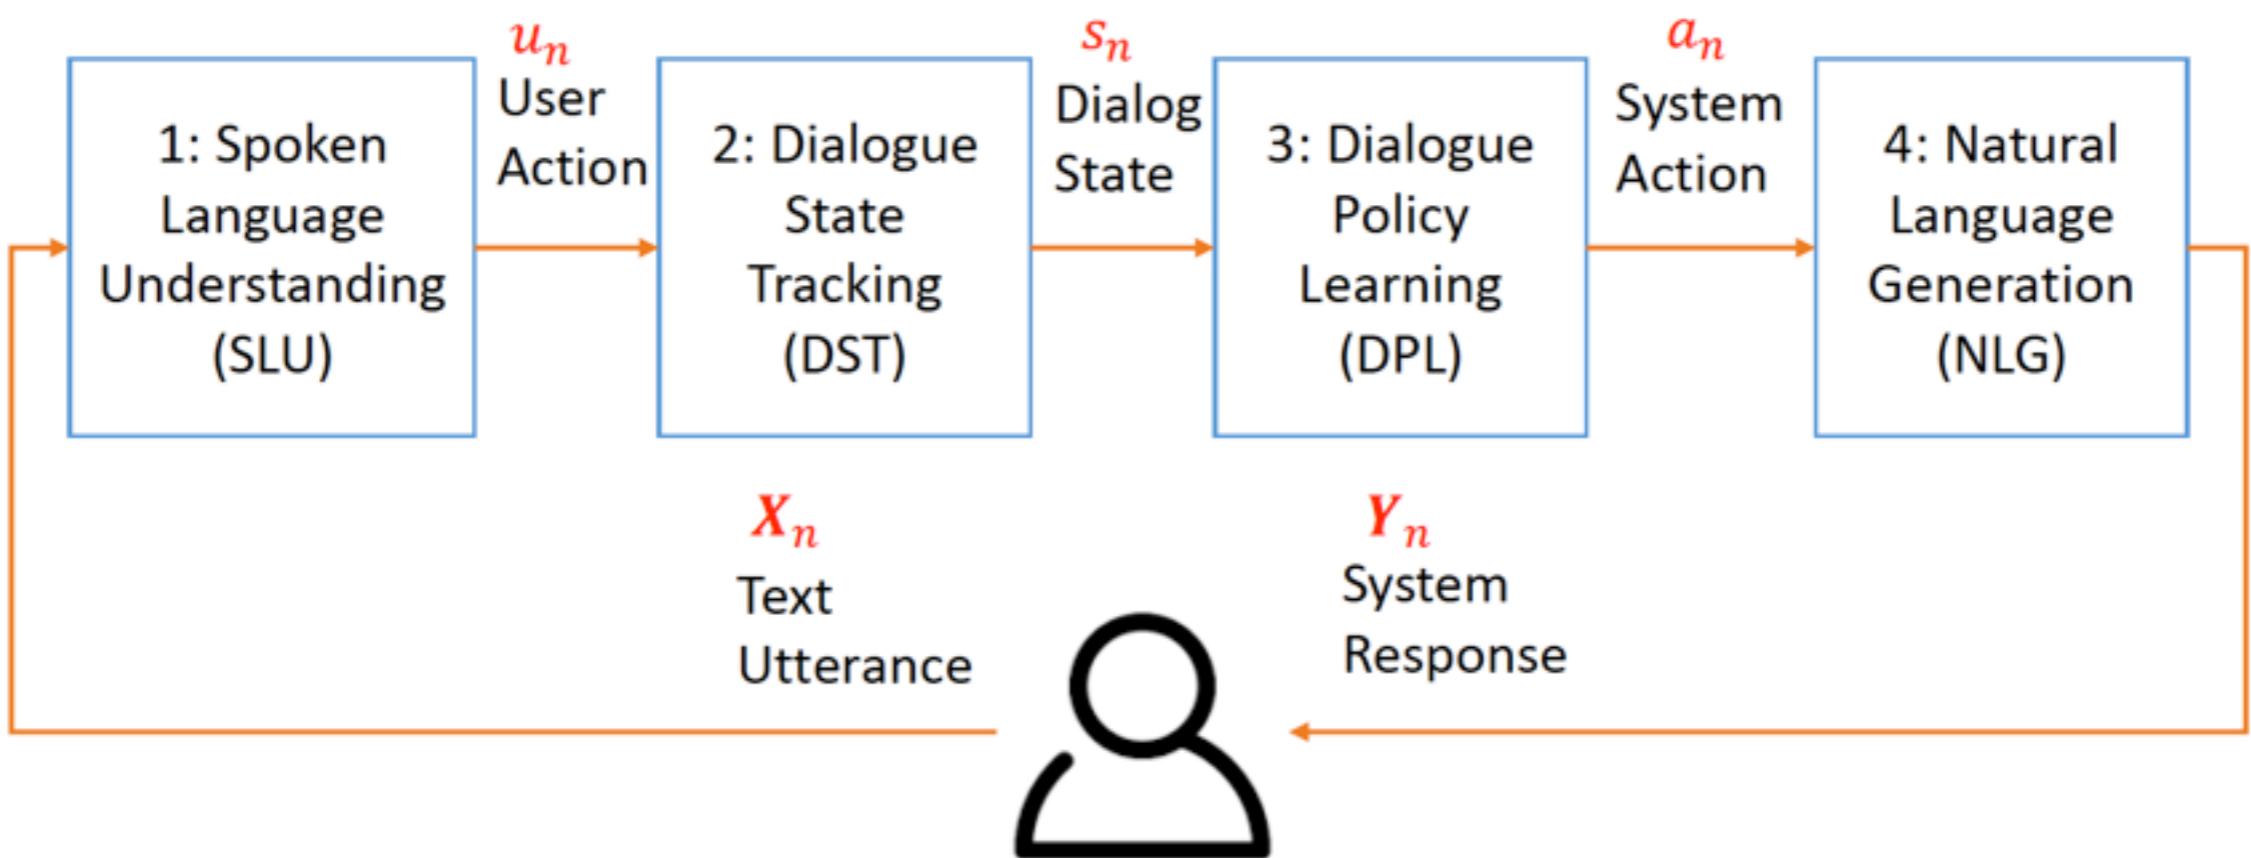


Task-oriented

Dialog System

Task-oriented Dialog System

- Modular: often 4 modules

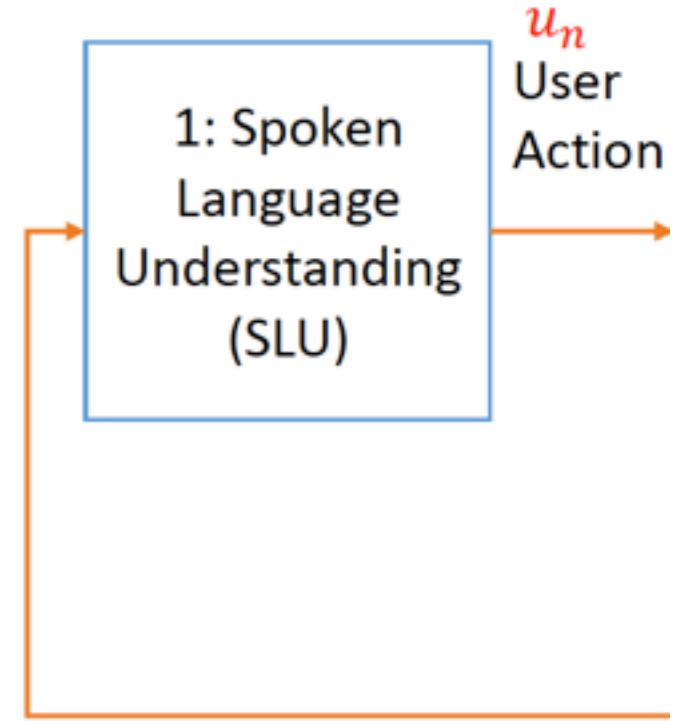


Task-oriented Dialog System

- Modular: often 4 modules
- Sub-modules, sub-tasks:
 - Spoken Language Understanding (SLU)
 - Dialogue State Tracking (DST)
 - Dialogue Policy Learning (DPL)
 - Natural Language Generation (NLG)

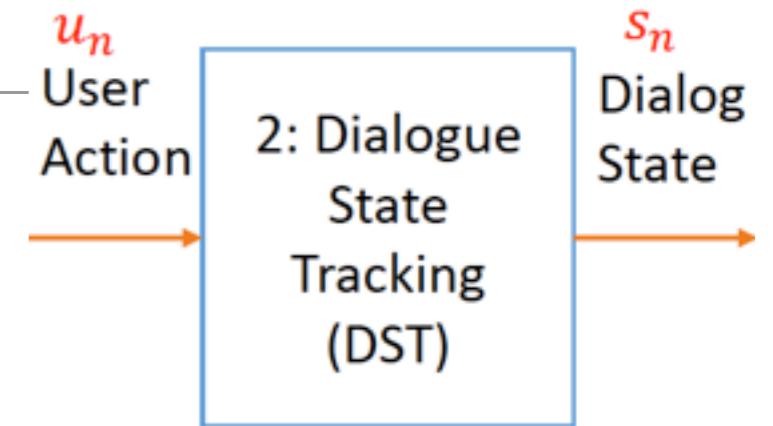
Task-oriented Dialog System

- Modular: often 4 modules
- Sub-modules, sub-tasks:
 - Spoken Language Understanding (SLU)
 - SLU turns natural language into user intention and slot-values, and it takes input and outputs structured user action
 - Dialogue State Tracking (DST)
 - Dialogue Policy Learning (DPL)
 - Natural Language Generation (NLG)



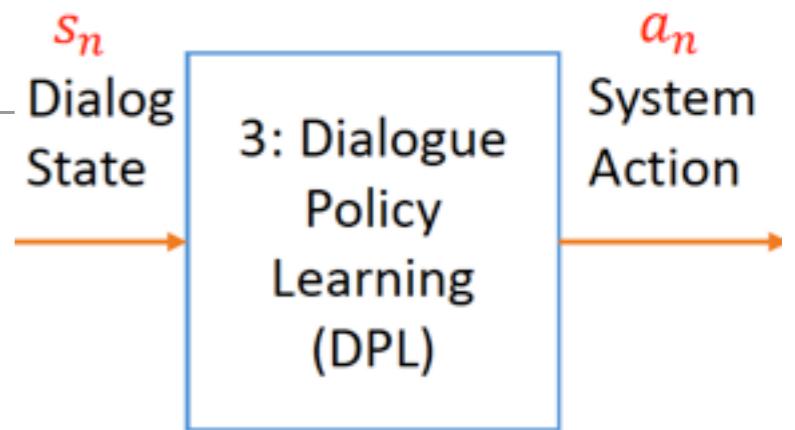
Task-oriented Dialog System

- Modular: often 4 modules
- Sub-modules, sub-tasks:
 - Spoken Language Understanding (SLU)
 - Dialogue State Tracking (DST)
 - DST tracks the current dialogue state, and outputs dialogue state
 - Dialogue Policy Learning (DPL)
 - Natural Language Generation (NLG)



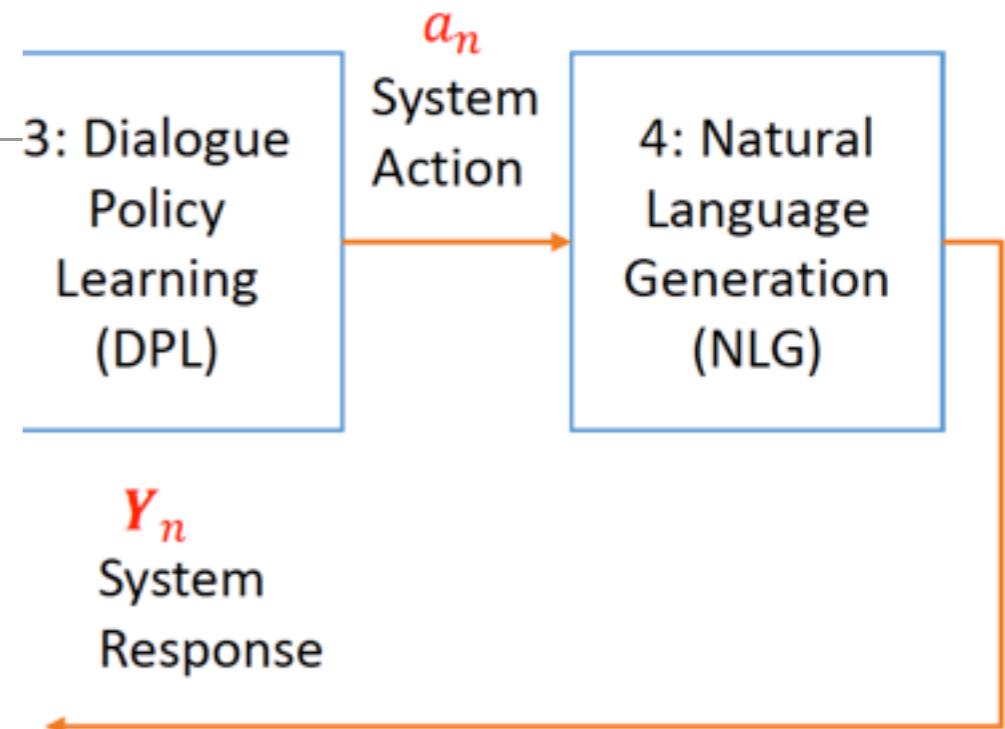
Task-oriented Dialog System

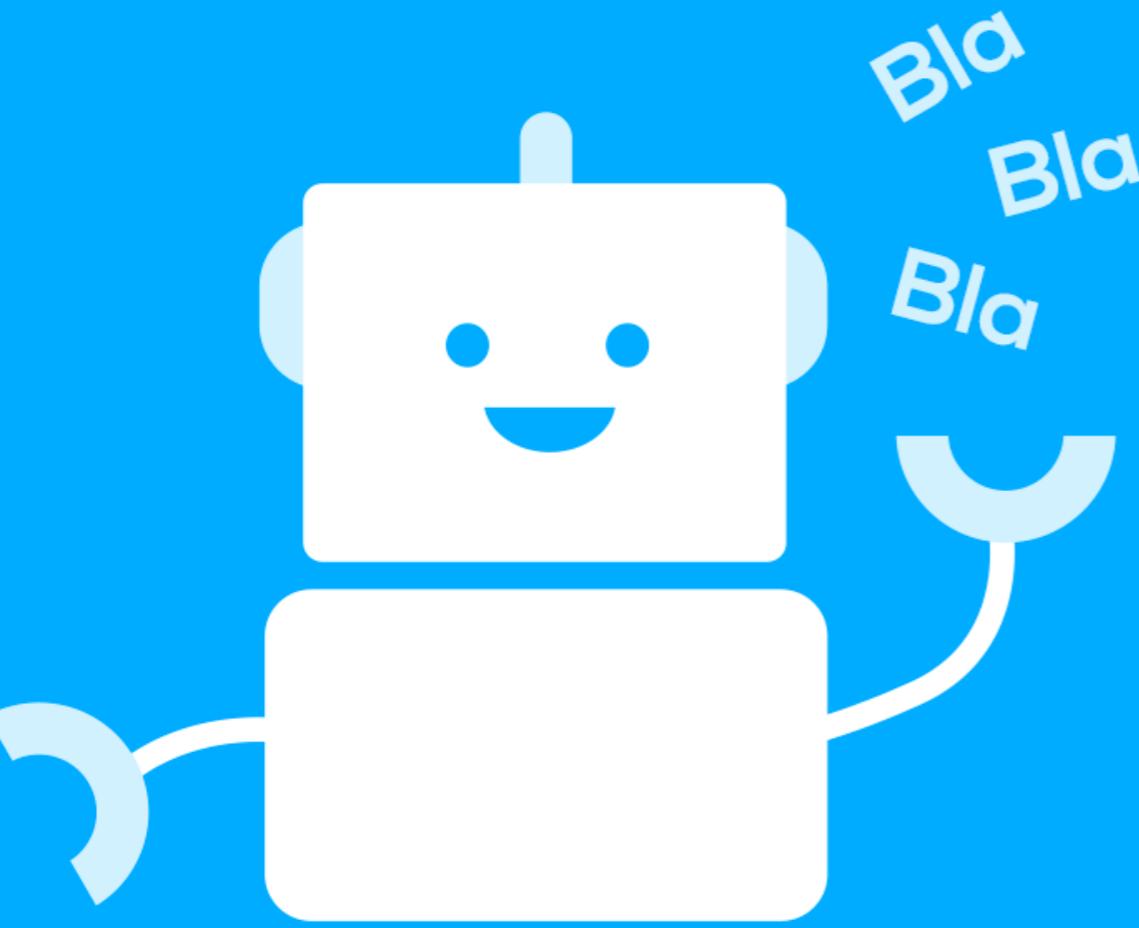
- Modular: often 4 modules
- Sub-modules, sub-tasks:
 - Spoken Language Understanding (SLU)
 - Dialogue State Tracking (DST)
 - Dialogue Policy Learning (DPL)
 - Policy decides which system action to take based on the dialogue state, and it takes dialogue state as input and outputs system action
 - Natural Language Generation (NLG)



Task-oriented Dialog System

- Modular: often 4 modules
- Sub-modules, sub-tasks:
 - Spoken Language Understanding (SLU)
 - Dialogue State Tracking (DST)
 - Dialogue Policy Learning (DPL)
 - Natural Language Generation (NLG)
 - NLG turns a system action into natural language, and it takes the system action as input and outputs the system response





Bla
Bla
Bla

vs.



Task-oriented

An Example

Example

- Wen et al., “A Network-based End-to-End Trainable Task-oriented Dialogue System”. EACL 2017.[7]
- <https://github.com/shawnwun/NNDIAL>

Task-oriented Neural Dialog System [7]

Can I have Korean

Little Seoul serves great Korean .

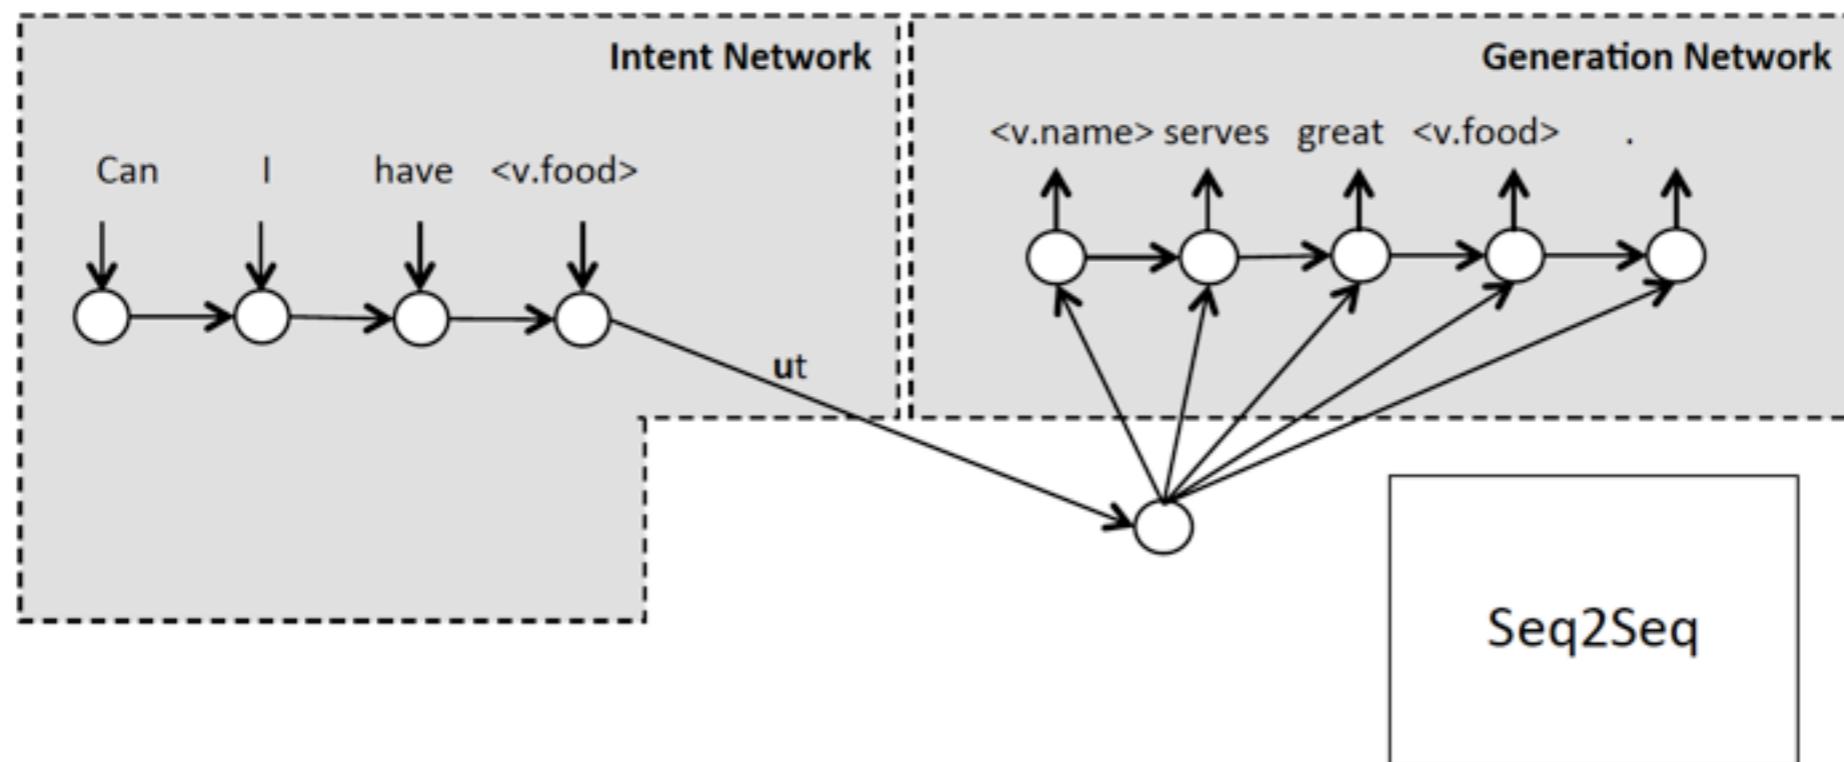
Task-oriented Neural Dialog System

Can I have <v.food>

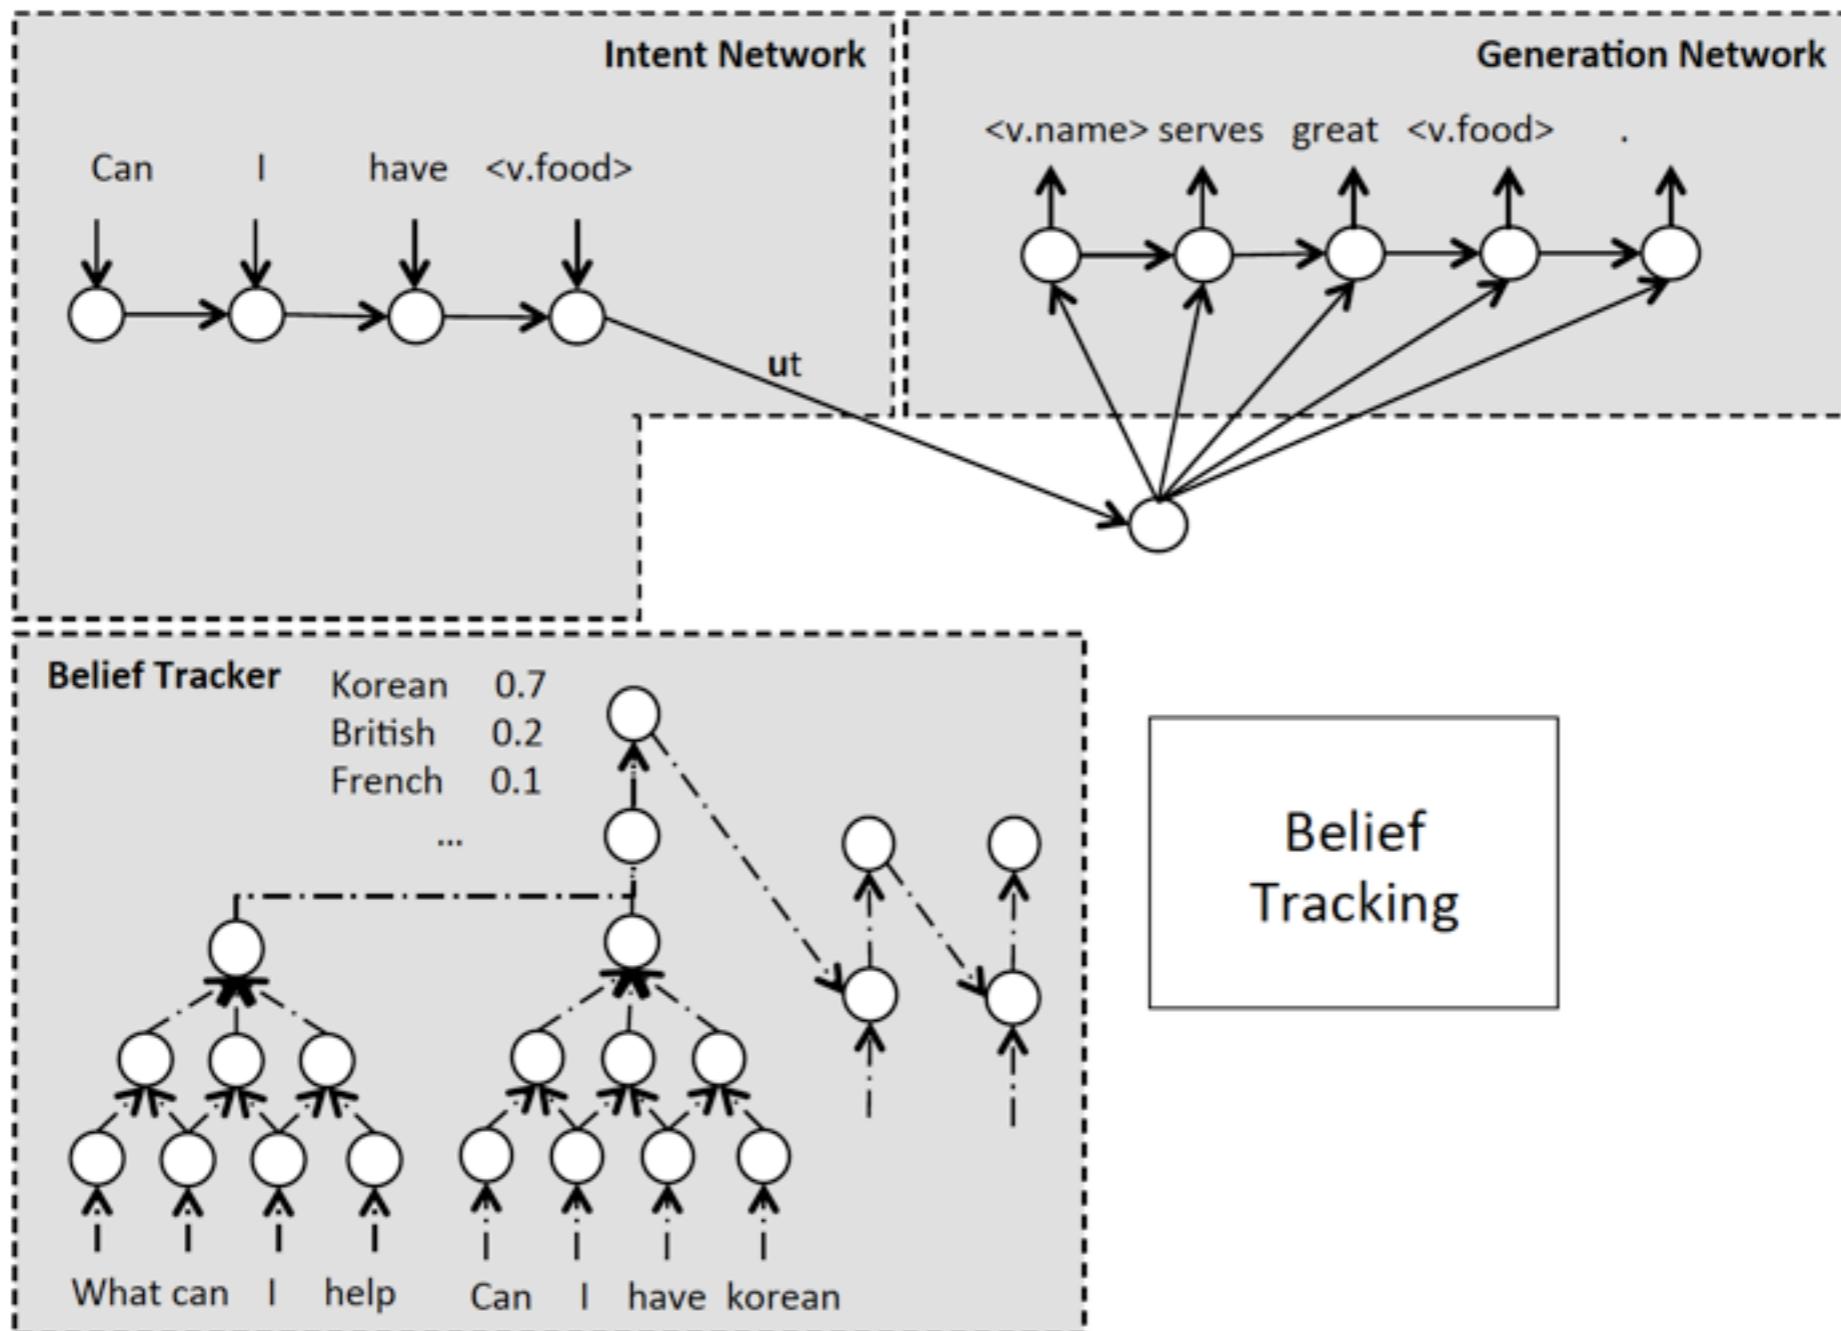
<v.name> serves great <v.food> .

Delexicalisation

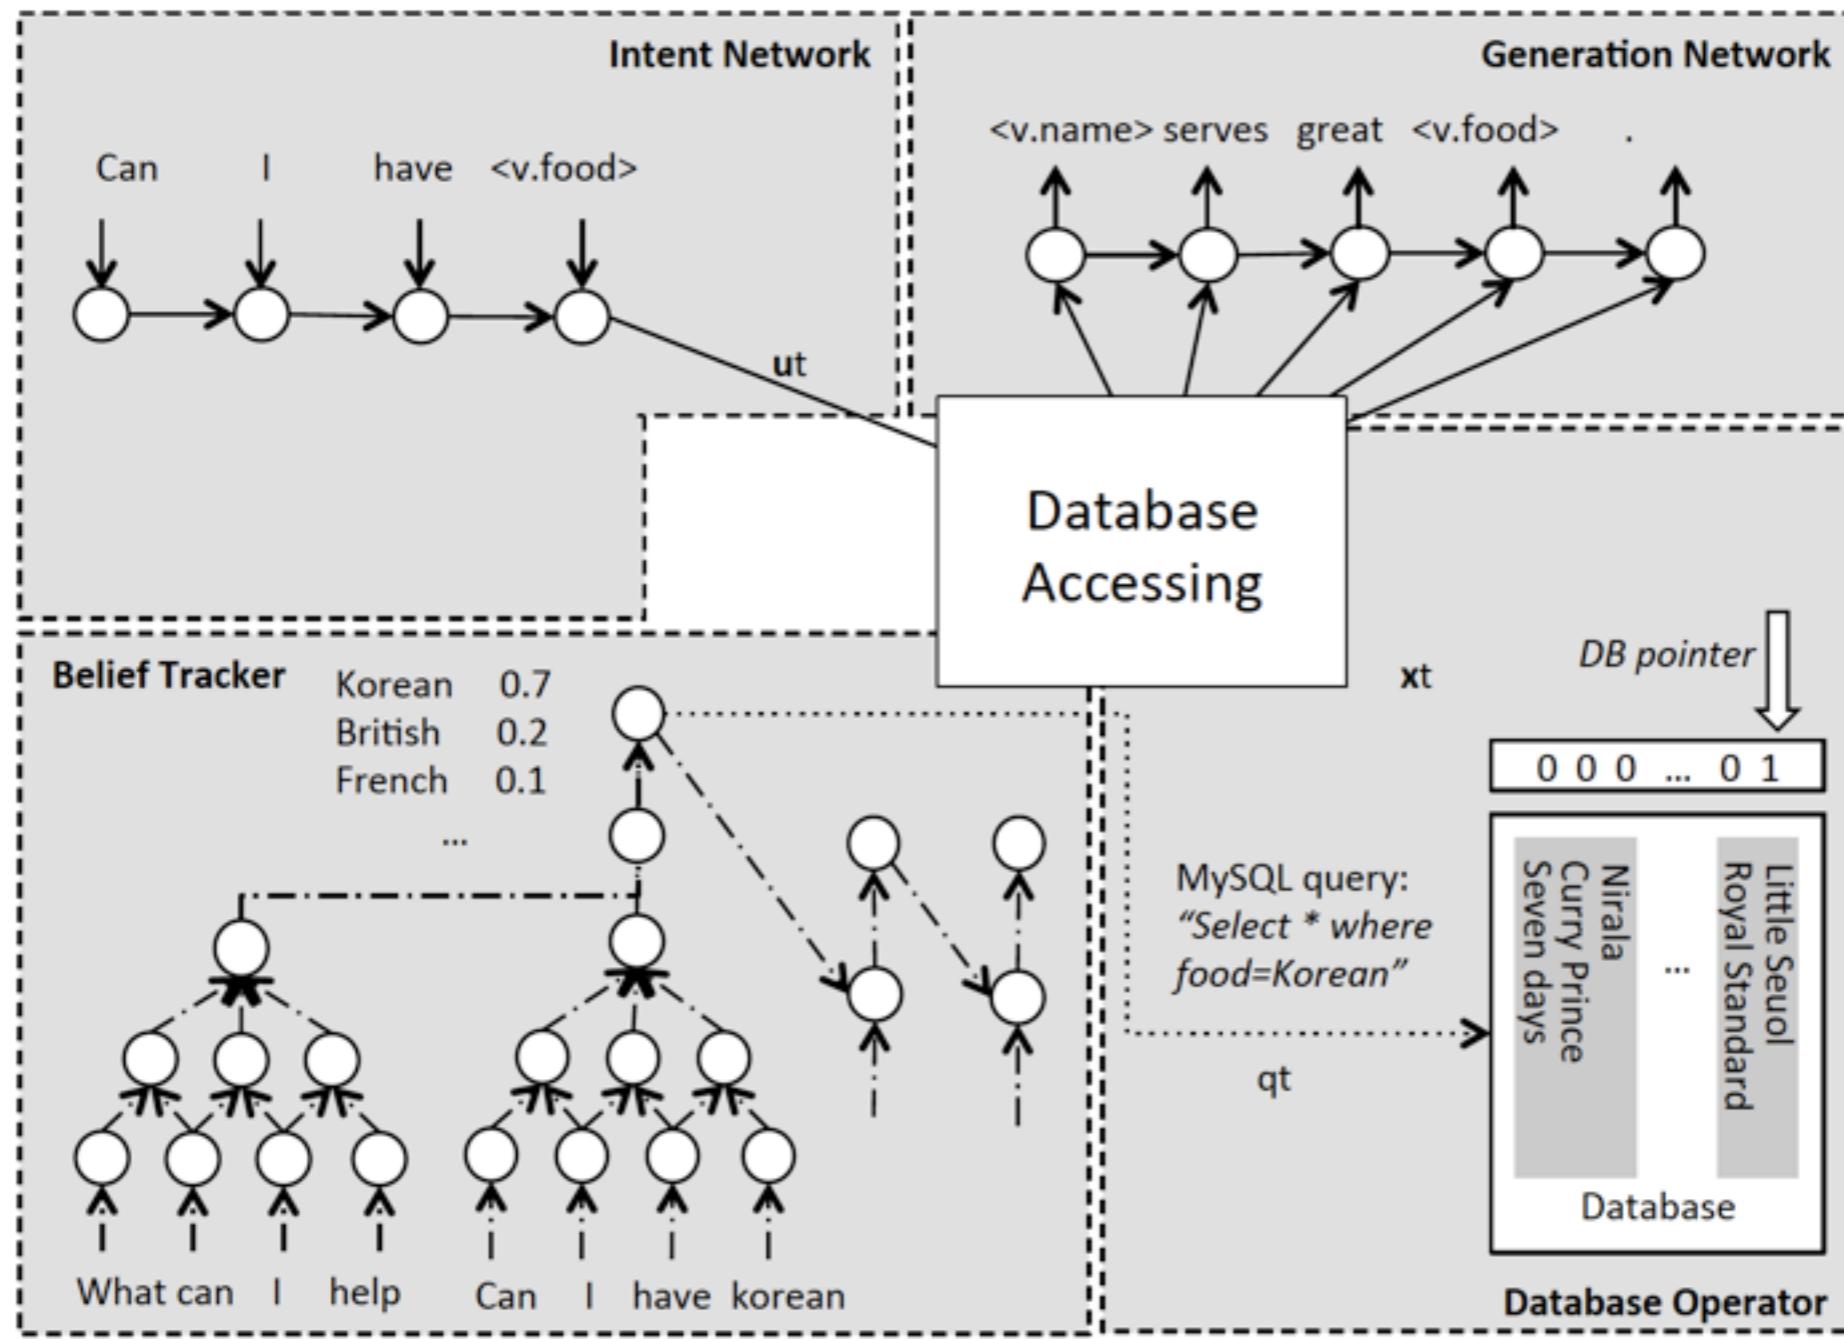
Task-oriented Neural Dialog System



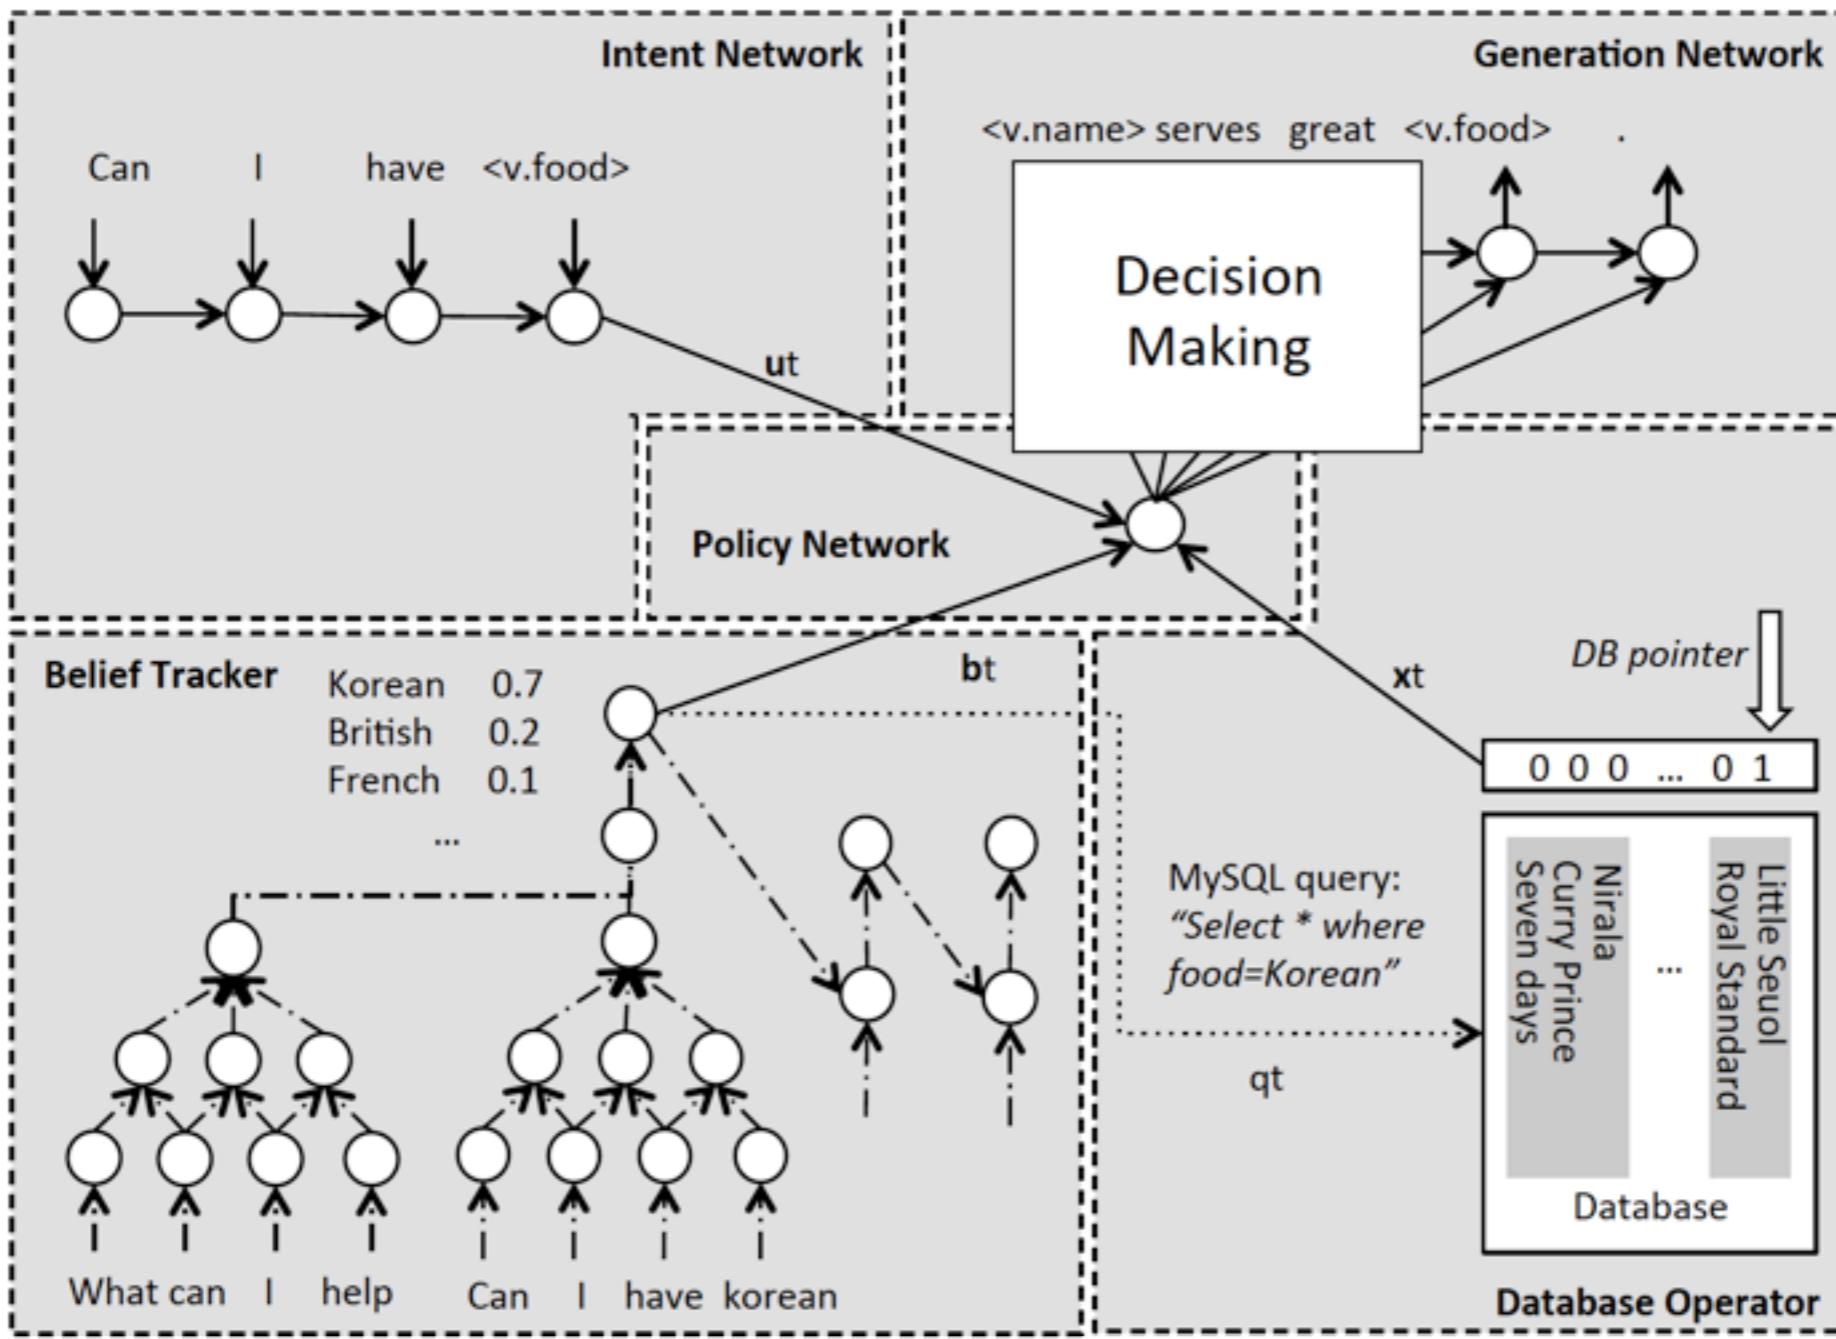
Task-oriented Neural Dialog System



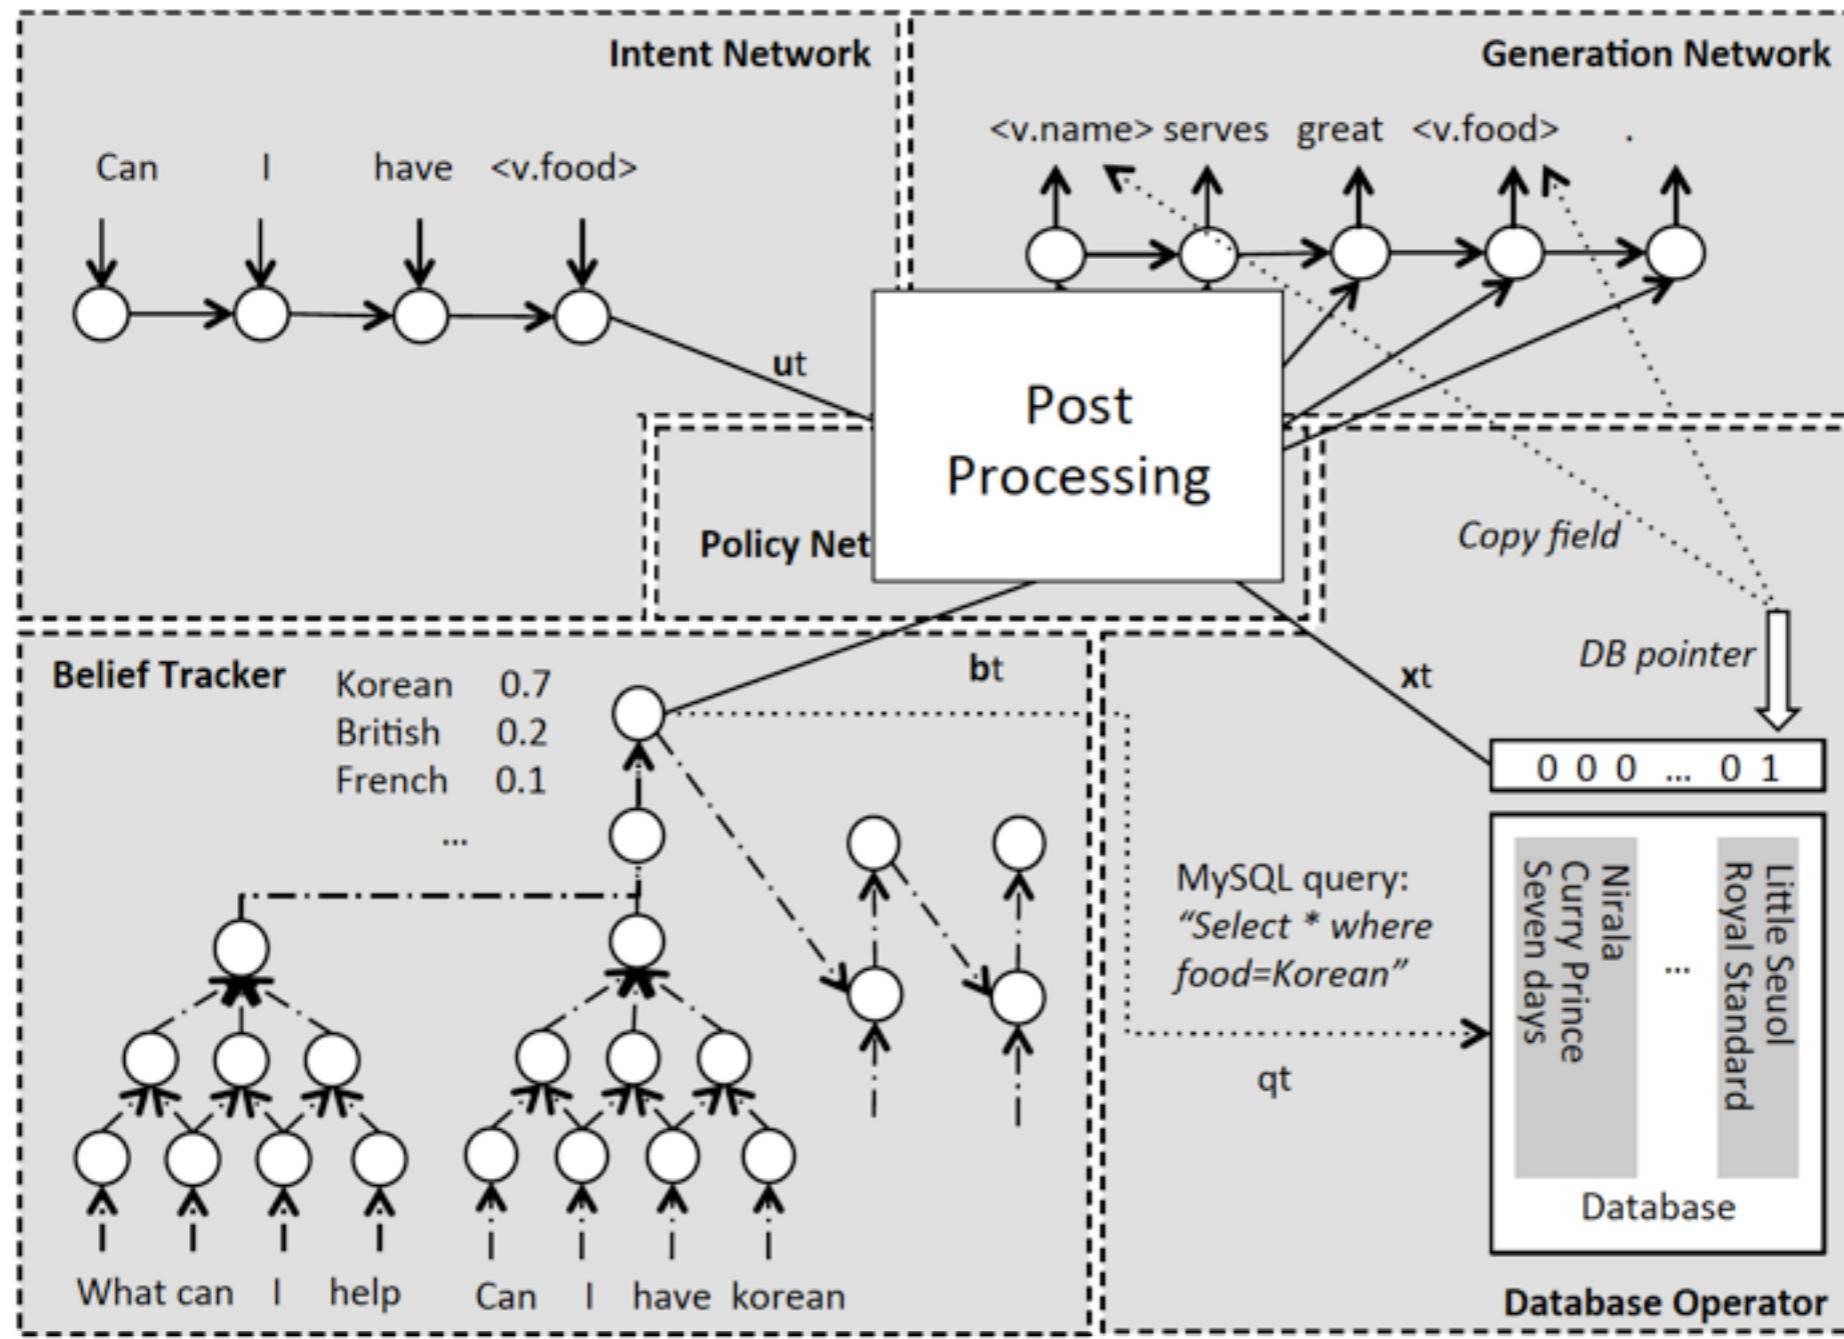
Task-oriented Neural Dialog System



Task-oriented Neural Dialog System



Task-oriented Neural Dialog System



Task-oriented Neural Dialog System [7]

Can I have Korean

Little Seoul serves great Korean .

Task-oriented Neural Dialog System

Can I have <v.food>

<v.name> serves great <v.food> .

Delexicalisation

Delexicalisation

- Replace slot values by **a symbol** representing its corresponding slot tokens. In the final postprocessing phase, these slot symbols are converted back to the corresponding slot values.
- A small knowledge graph defining the slot-value pairs the system can talk about for a particular task.
- Features and corpus can be found in [12].
- *In this paper, the system was designed to assist users to find a restaurant in the Cambridge, UK area. There are three informable slots (food, pricerange, area) that users can use to constrain the search and six requestable slots (address, phone, postcode plus the three informable slots) that the user can ask a value for once a restaurant has been offered. There are 99 restaurants in the DB. [7]*

Delexicalisation

- Replace slot values by **a symbol** representing its corresponding slot tokens. In the final postprocessing phase, these slot symbols are converted back to the corresponding slot values.
- A small knowledge graph defining the slot-value pairs the system can talk about for a particular task.
- Features and corpus can be found in [12].

| Delexicalised token | Examples |
|-------------------------|---------------------------|
| informable slot token | <s.food>, <s.area>,... |
| informable value token | <v.food>, <v.area>,... |
| requestable slot token | <s.phone>,<s.address>,... |
| requestable value token | <v.phone>,<v.address>,... |

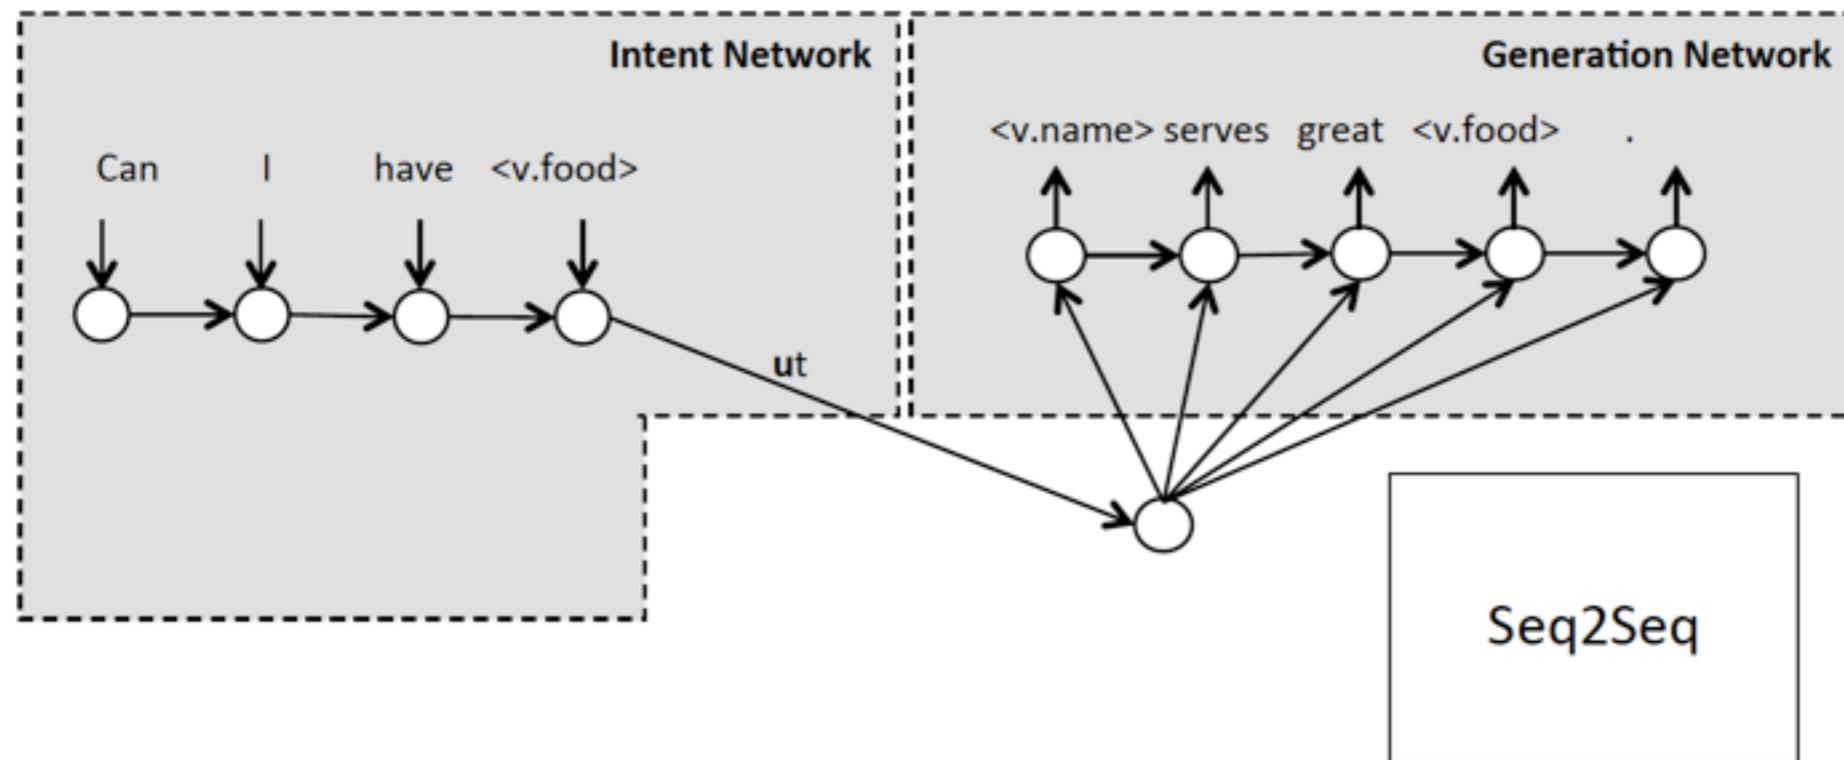
Delexicalisation

```
549     def prepareSlotValues(self):
550
551         print '\tprepare slot value templates ...'
552         # put db requestable values into s2v
553         for e in self.db:
554             for s,v in e.iteritems():
555                 if self.s2v['requestable'].has_key(s):
556                     self.s2v['requestable'][s].append(v.lower())
557                 if self.s2v['other'].has_key(s):
558                     self.s2v['other'][s].append(v.lower())
559         # sort values
560         for s,vs in self.s2v['informable'].iteritems():
561             self.s2v['informable'][s] = sorted(list(set(vs)))
562         for s,vs in self.s2v['requestable'].iteritems():
563             self.s2v['requestable'][s] = sorted(list(set(vs)))
564         for s,vs in self.s2v['other'].iteritems():
565             self.s2v['other'][s] = sorted(list(set(vs)))
566
567         # make a 1-on-1 mapping for delexicalisation
568         self.supervalues = []
569         self.values = []
570         self.slots = []
```

Delexicalisation

```
519     def delexicalise(self,utt,mode='all'):
520         inftoks =  ['[VALUE_+' + s.upper() + ']' for s in self.s2v['informable'].keys()] + \
521                     ['[SLOT_ ' + s.upper() + ']' for s in self.s2v['informable'].keys()] + \
522                     ['[VALUE_DONTCARE]', '[VALUE_NAME]' ] + \
523                     ['[SLOT_ ' + s.upper() + ']' for s in self.s2v['requestable'].keys()]
524         reqtoks =  ['[VALUE_+' + s.upper() + ']' for s in self.s2v['requestable'].keys()]
525         for i in range(len(self.values)):
526             # informative mode, preserving location information
527             if mode=='informative' and self.slots[i] in inftoks:
528                 tok = self.slots[i]+':'+(self.supervalues[i]).replace(' ','-')
529                 utt = (' '+utt+' ').replace(' '+self.values[i]+' ',' '+tok+' ')
530                 utt = utt[1:-1]
531             # requestable mode
532             elif mode=='requestable' and self.slots[i] in reqtoks:
533                 utt = (' '+utt+' ').replace(' '+self.values[i]+' ',' '+self.slots[i]+' ')
534                 utt = utt[1:-1]
535             elif mode=='all':
536                 tok = self.slots[i]+':'+(self.supervalues[i]).replace(' ','-') \
537                         if self.slots[i] in inftoks else self.slots[i]
538                 utt = (' '+utt+' ').replace(' '+self.values[i]+' ',' '+tok+' ')
539                 utt = utt[1:-1]
540         utt = re.sub(digitpat,'[VALUE_COUNT]',utt)
541         return utt
```

Task-oriented Neural Dialog System



Intent Network

- convert user input (utterances) into distributed representation at every speaker turn.
- can be viewed as the encoder in the sequence-to-sequence learning framework, e.g. LSTM or CNN. The example code repository adopts LSTM.

* The encoder modules contain:

- LSTM encoder : an LSTM network that encodes the user utterance.
- RNN+CNN tracker : a set of slot trackers that keep track of each slot/value pair across turns.
- DB operator : a discrete database accessing component.

* The decoder modules contain:

- Policy network : a decision-making module that produces the conditional vector for decoding.
- LSTM decoder : an LSTM network that generates the system response.

Recall Long-Short Term Memory (LSTM)

- ◆ Three sigmoid **gates** and One **cell**



$$\mathbf{i}_t = \sigma(\mathbf{W}_{wi}\mathbf{x}_t + \mathbf{W}_{hi}\mathbf{h}_{t-1})$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \hat{\mathbf{c}}_t$$

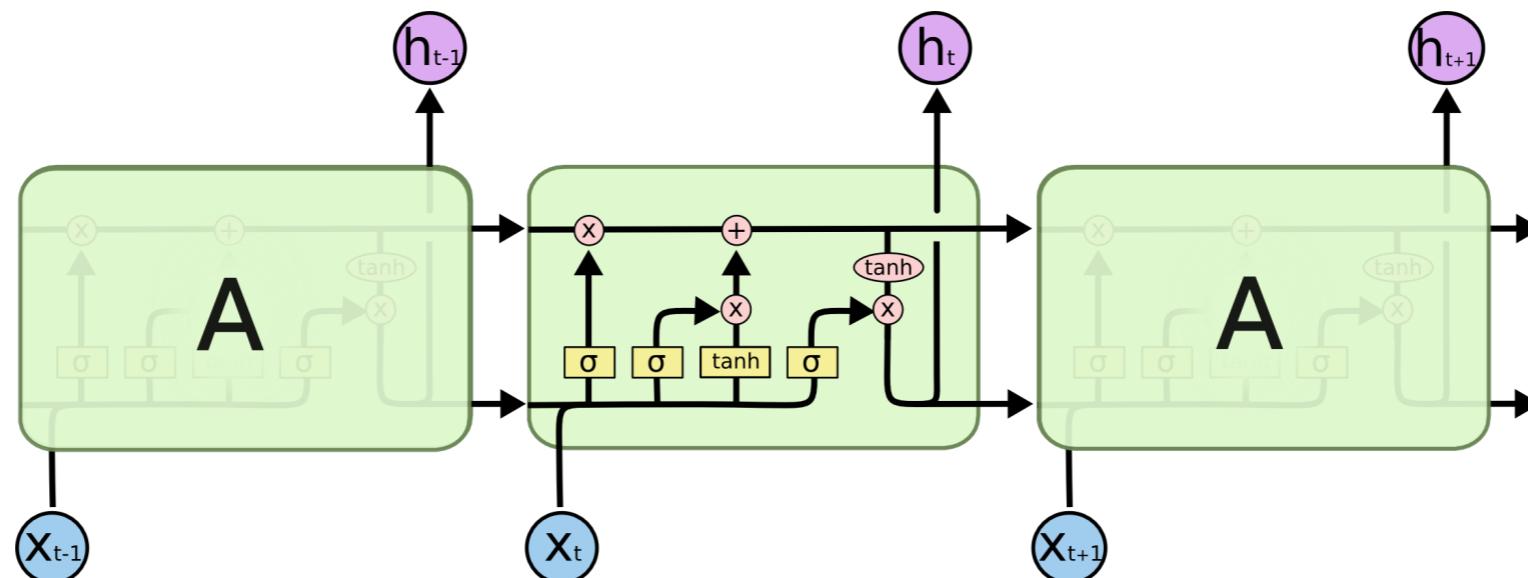
$$\mathbf{f}_t = \sigma(\mathbf{W}_{wf}\mathbf{x}_t + \mathbf{W}_{hf}\mathbf{h}_{t-1})$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_{wo}\mathbf{x}_t + \mathbf{W}_{ho}\mathbf{h}_{t-1})$$

$$\hat{\mathbf{c}}_t = \tanh(\mathbf{W}_{wc}\mathbf{x}_t + \mathbf{W}_{hc}\mathbf{h}_{t-1})$$

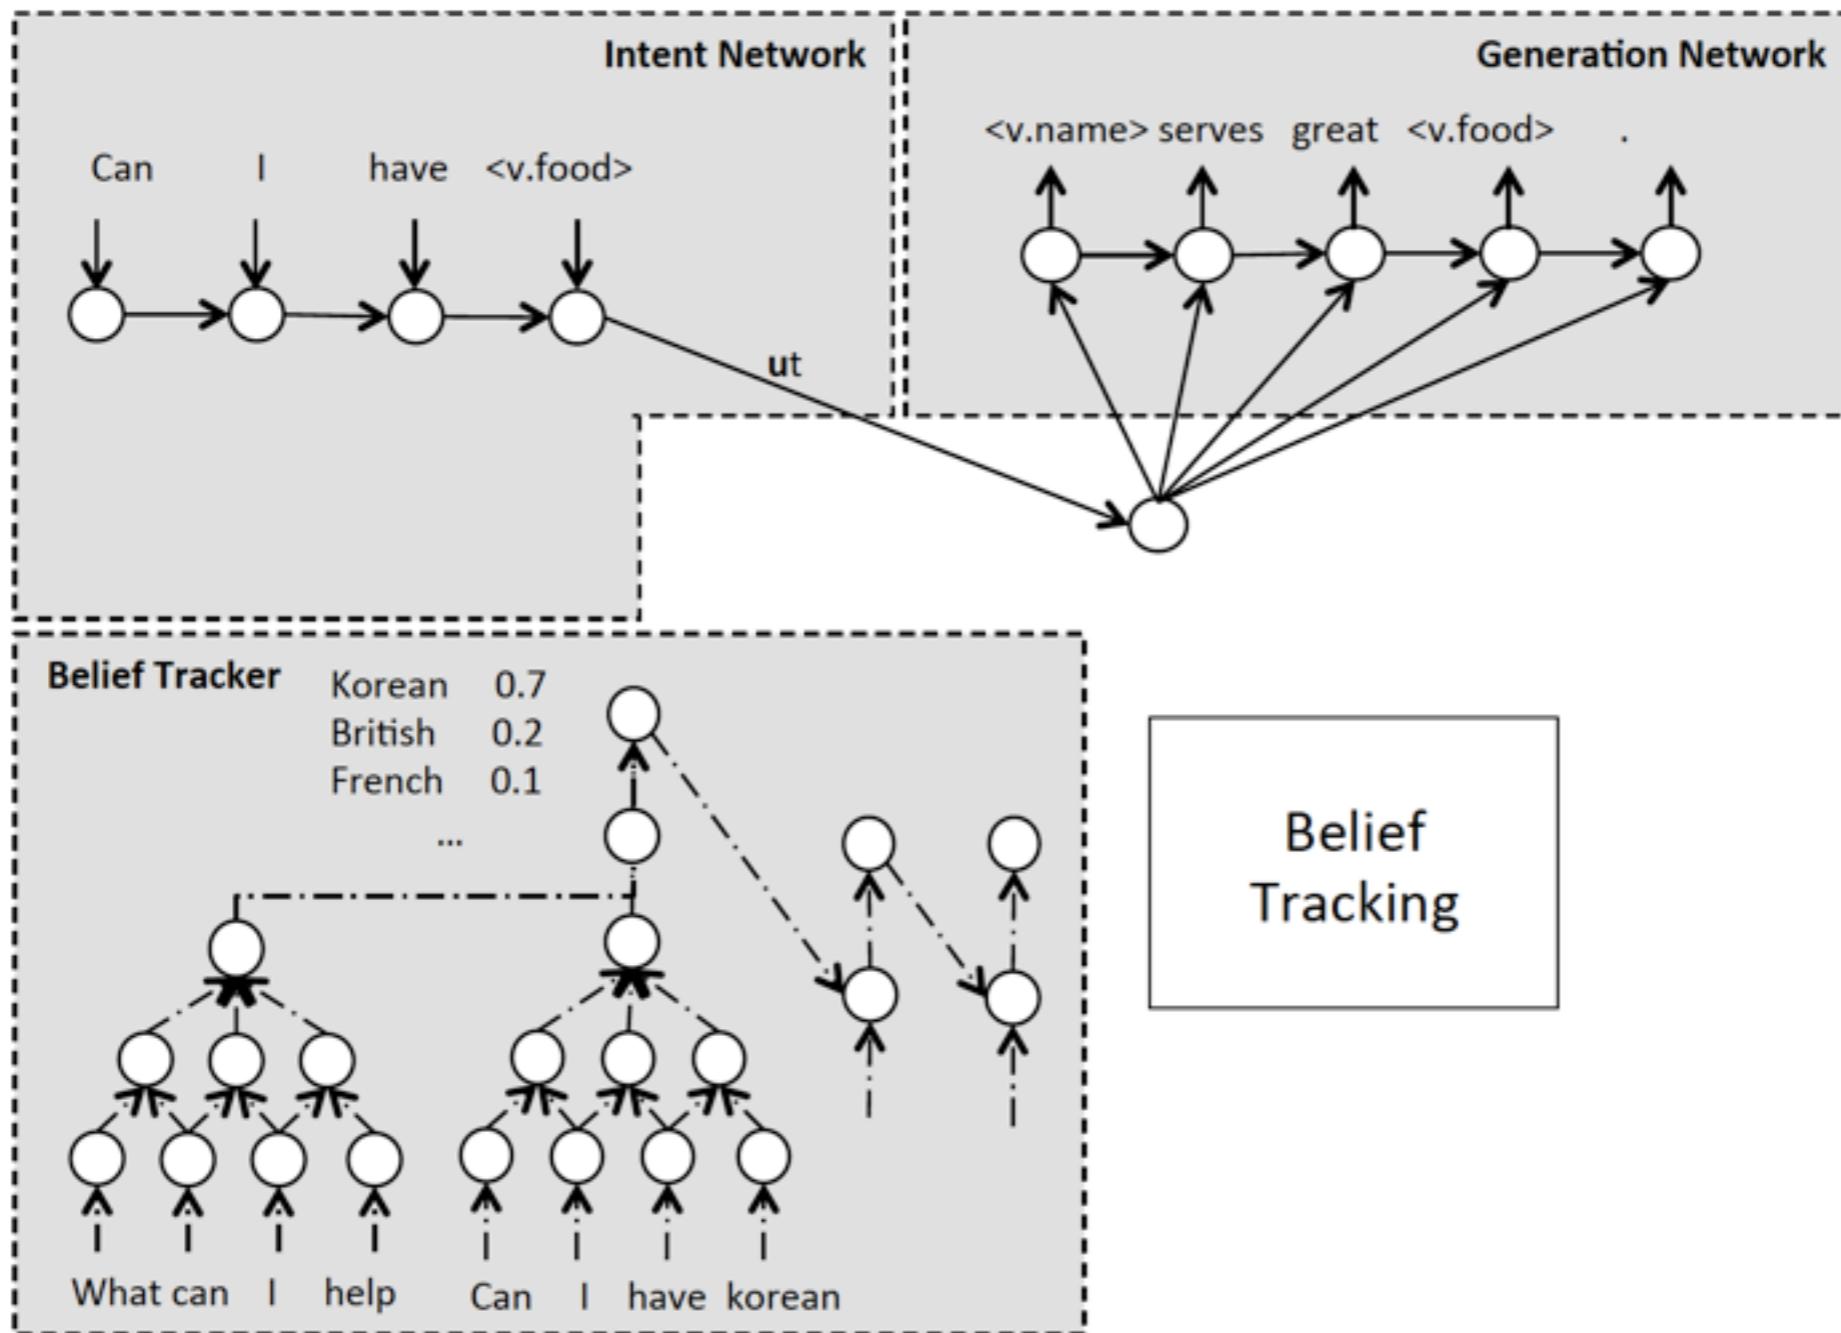
- ◆ LSTM



Intent Network

```
164 # LSTM Encoder
165 class LSTMEncoder(BaseNNModule):
166
167     def __init__(self, vocab_size, ihidden_size):
168
169         # parameters for encoder bilSTM
170         self.dih = ihidden_size
171         self.di = vocab_size
172         self.iWgate = theano.shared(0.3 * np.random.uniform(-0.1,
173                                         (self.dih, self.dih*4)).astype(theano.config.floatX))
174         self.iUgate = theano.shared(0.3 * np.random.uniform(-0.1,
175                                         (self.dih, self.dih*4)).astype(theano.config.floatX))
176         self.Wemb = theano.shared(0.3 * np.random.uniform(-0.1,
177                                         (self.di, self.dih)).astype(theano.config.floatX))
178         self.bf = theano.shared(2.0*np.ones((ihidden_size, 1)).astype(theano.config.floatX))
179
180
181     self.params = [
182         self.iWgate, self.iUgate, self.Wemb, self.bf ]
183
184     # initial states
185     self.ih0 = theano.shared(np.zeros((self.dih),\n                                         dtype=theano.config.floatX))
186     self.ic0 = theano.shared(np.zeros((self.dih),\n                                         dtype=theano.config.floatX))
187
188
251     # Bidirectional encoding
252     def bidirectional_encode(fEncoder, bEncoder, sent, leng):
253
254         fw_intent_t = fEncoder.encode(sent, leng)
255         bw_intent_t = bEncoder.encode(T.concatenate([
256             sent[:leng][::-1], sent[leng:]], axis=0), leng)
257
258         intent_t = T.concatenate([
259             fw_intent_t, bw_intent_t], axis=0)
260
261     return intent_t
262
263
264     def bidirectional_read(fEncoder, bEncoder, sent):
265
266         fw_intent_t = fEncoder.read(sent)
267         bw_intent_t = bEncoder.read(sent[::-1])
268
269         intent_t = np.concatenate([fw_intent_t, bw_intent_t], axis=0)
270
271     return intent_t
```

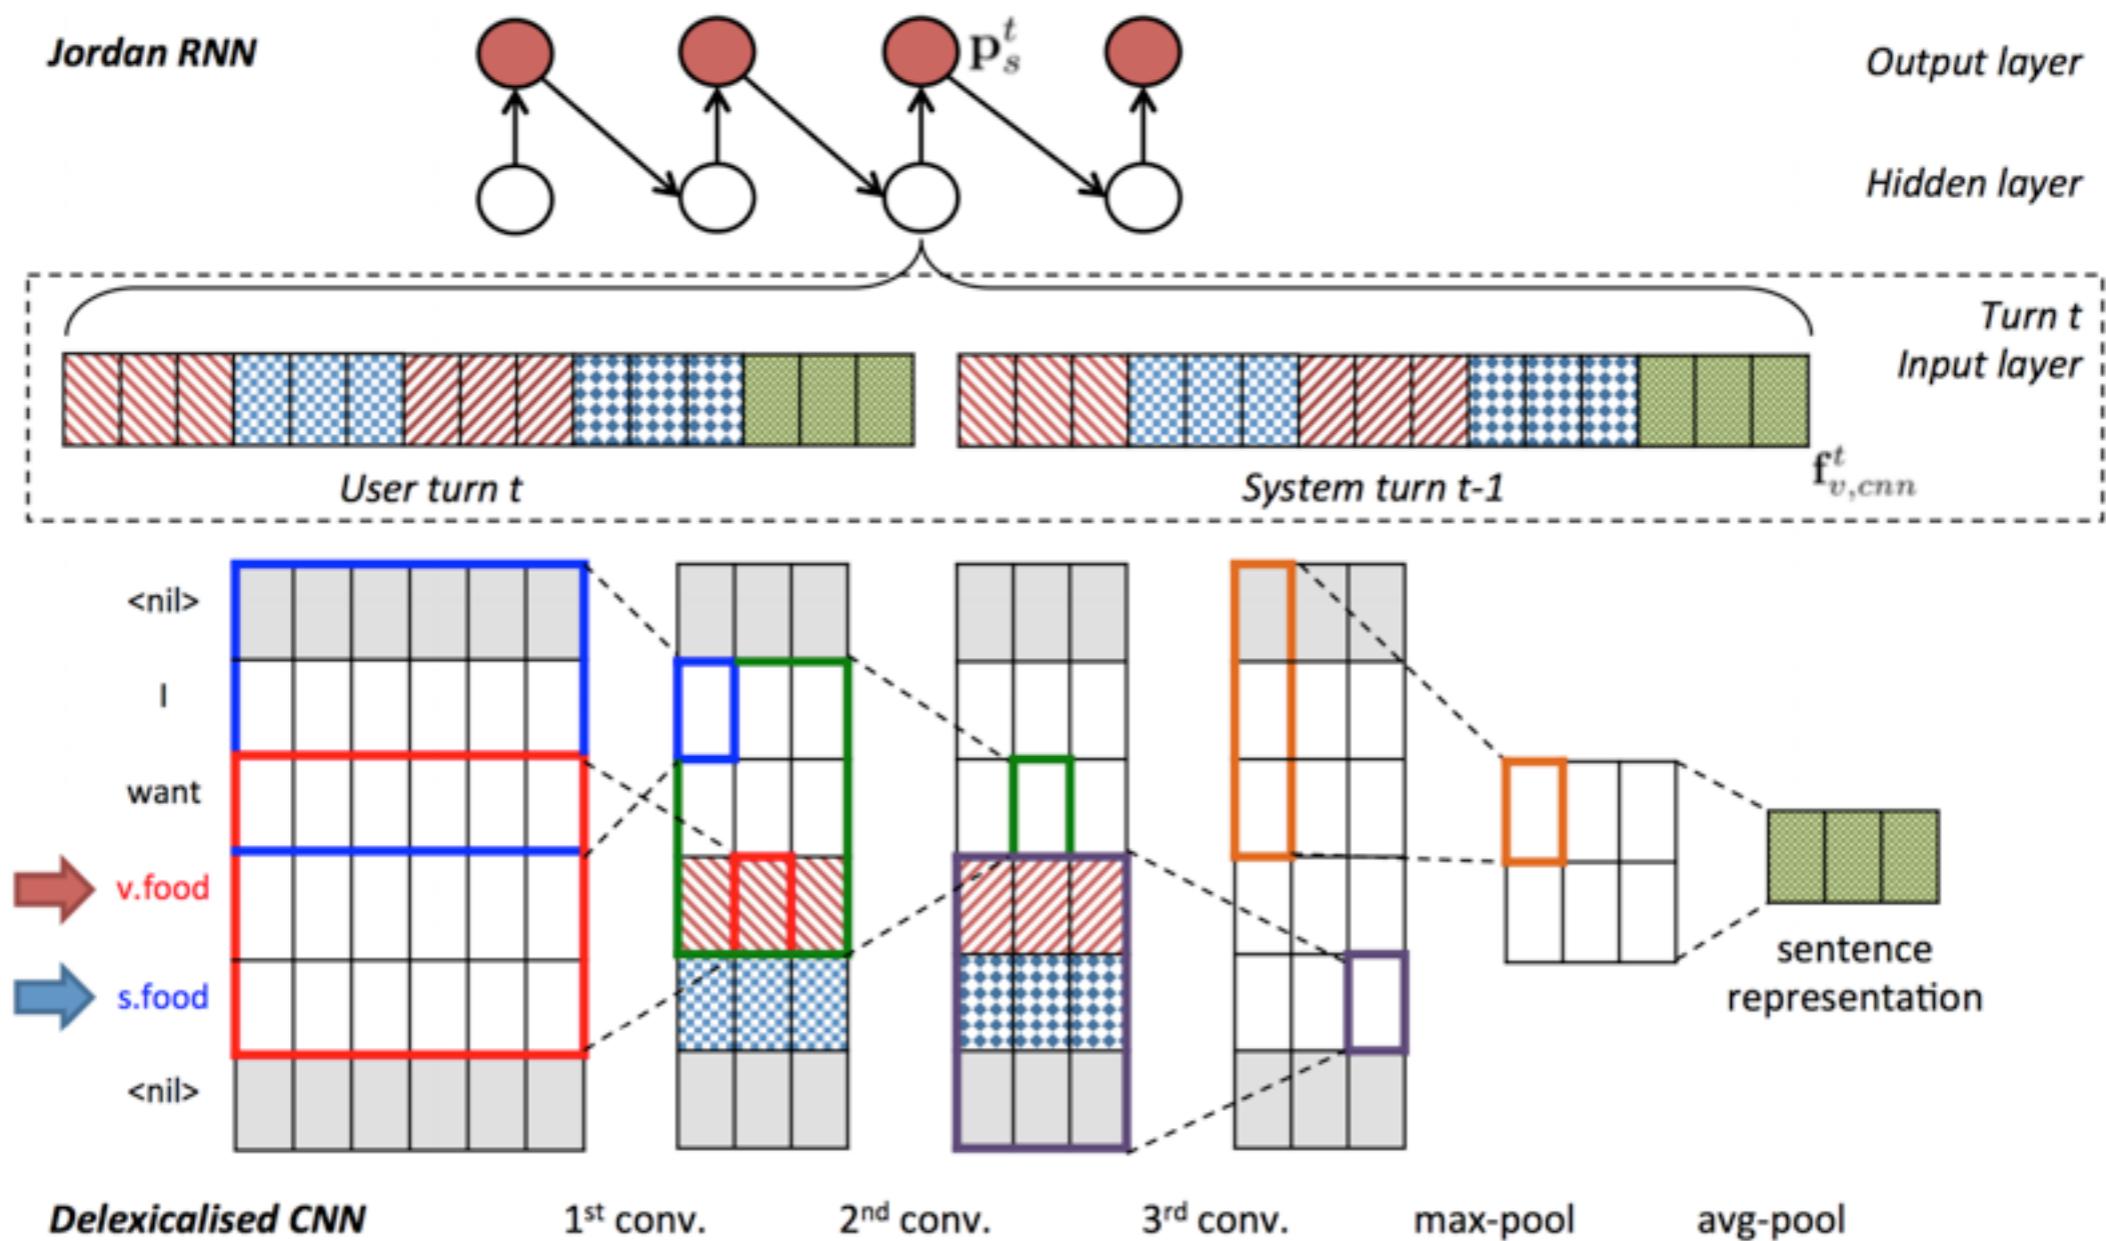
Task-oriented Neural Dialog System



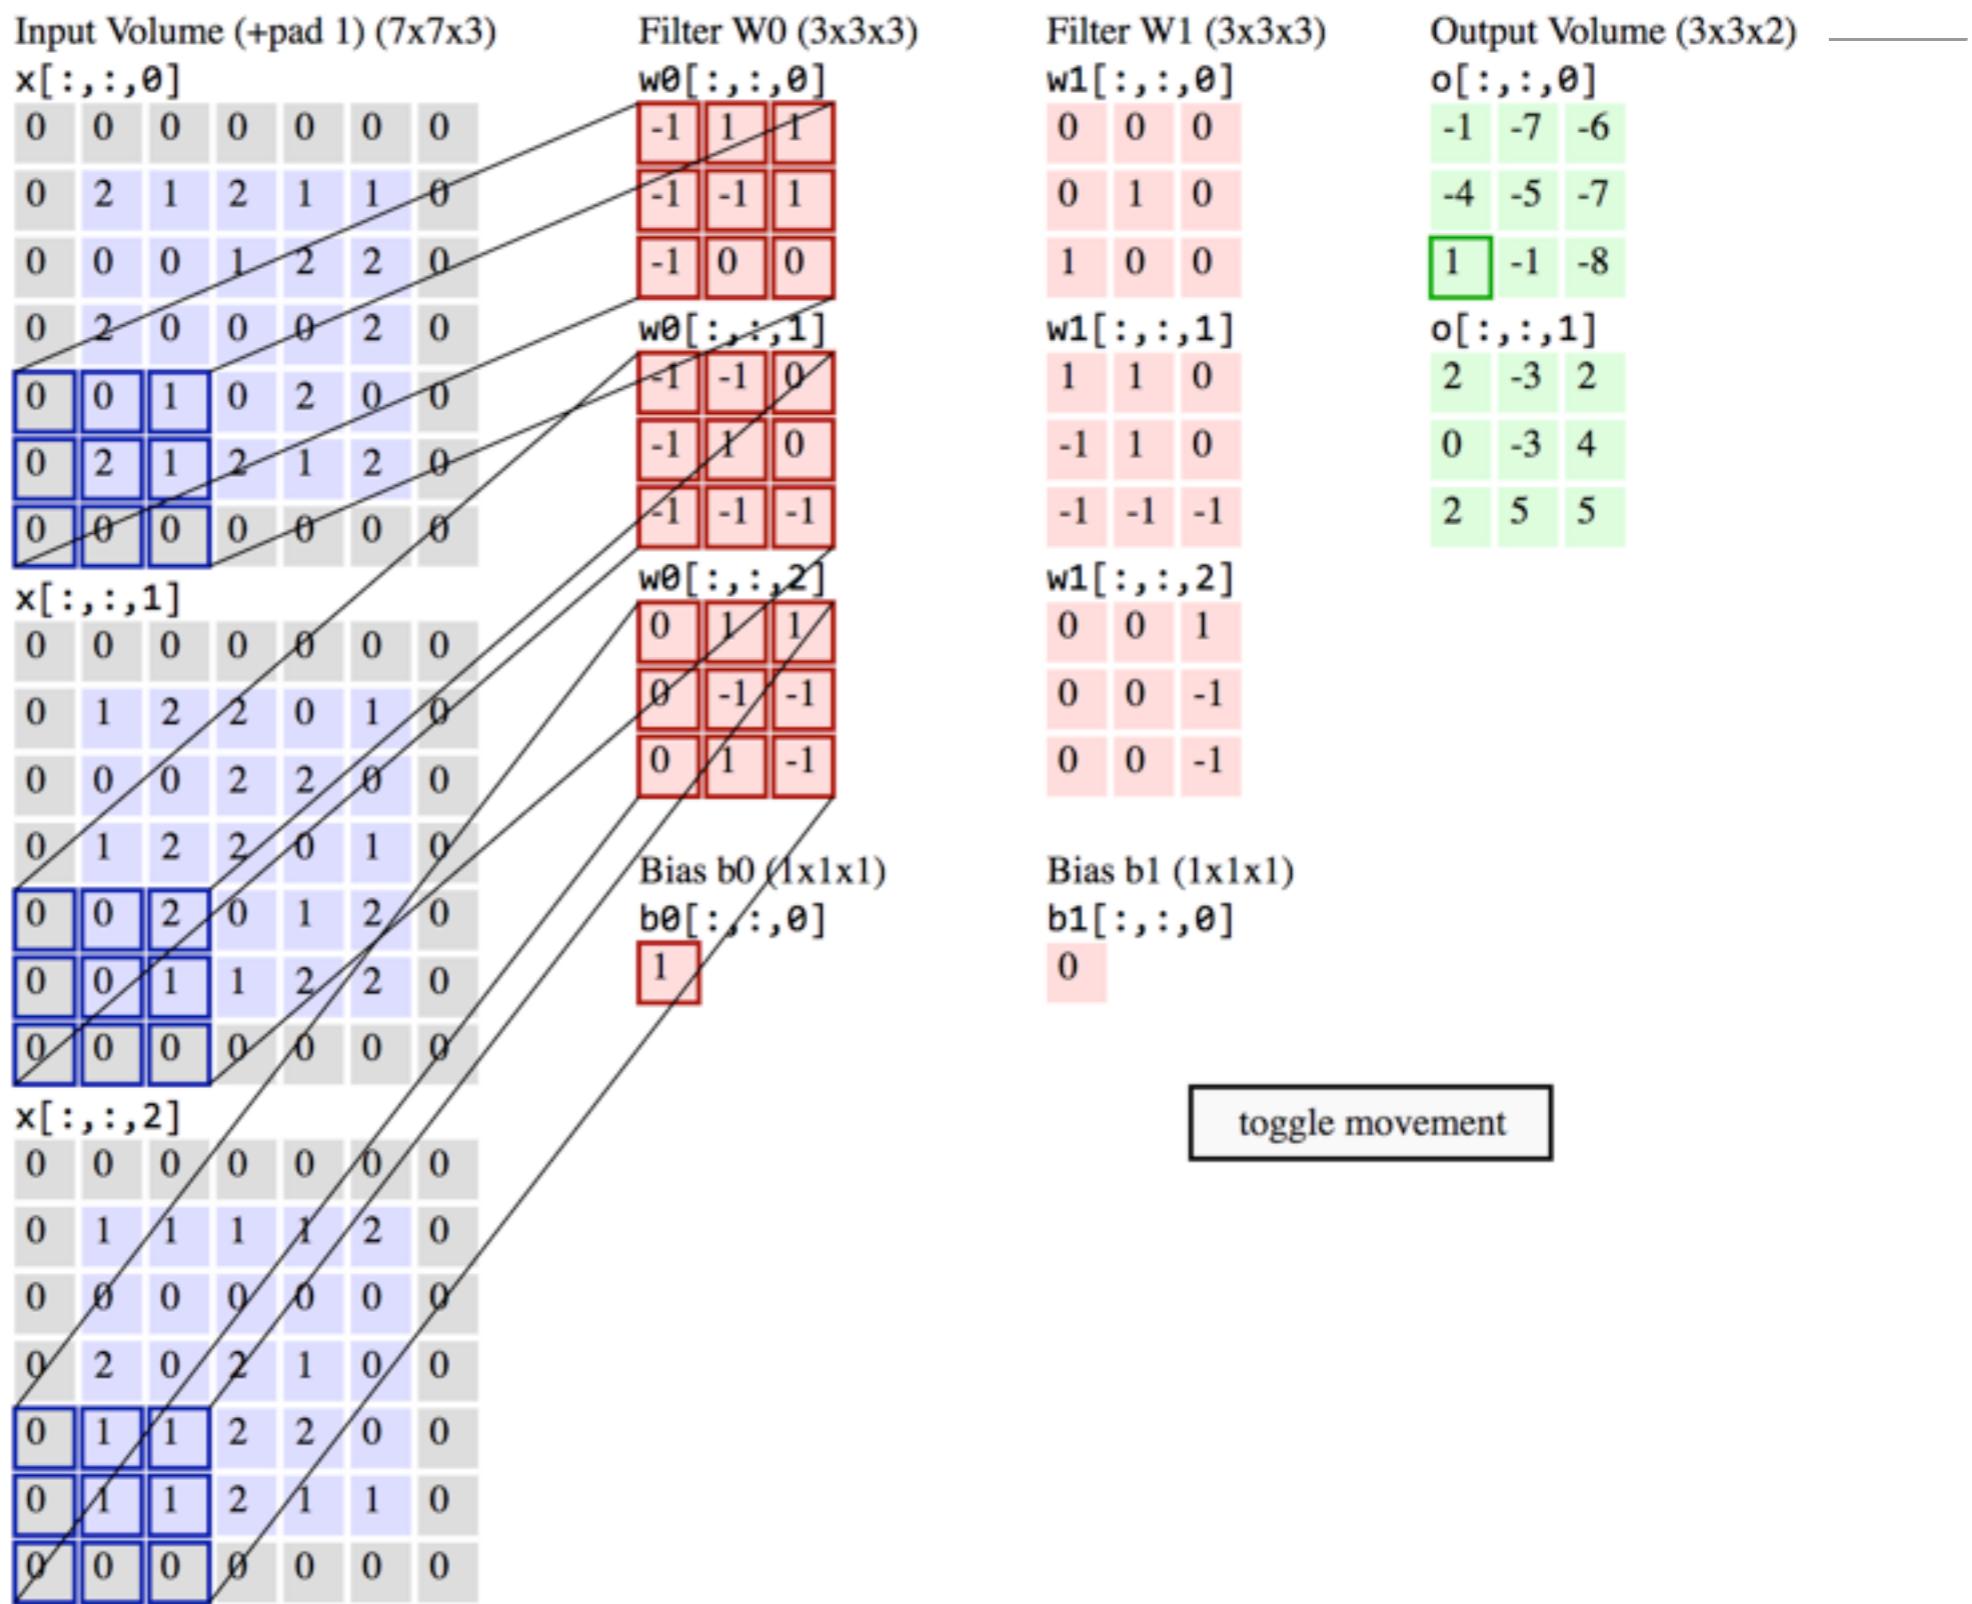
Belief Tracker

- convert user input (utterances) into a probability distribution over slot-value pairs, which is called the belief states, a multinomial distribution for **each** informable slot, and a binary distribution for **each** requestable slot.
- a set of belief trackers over a set of slot-value pairs.
- Belief tracking (also called Dialogue State tracking) provides the core of a task-oriented spoken dialogue system.

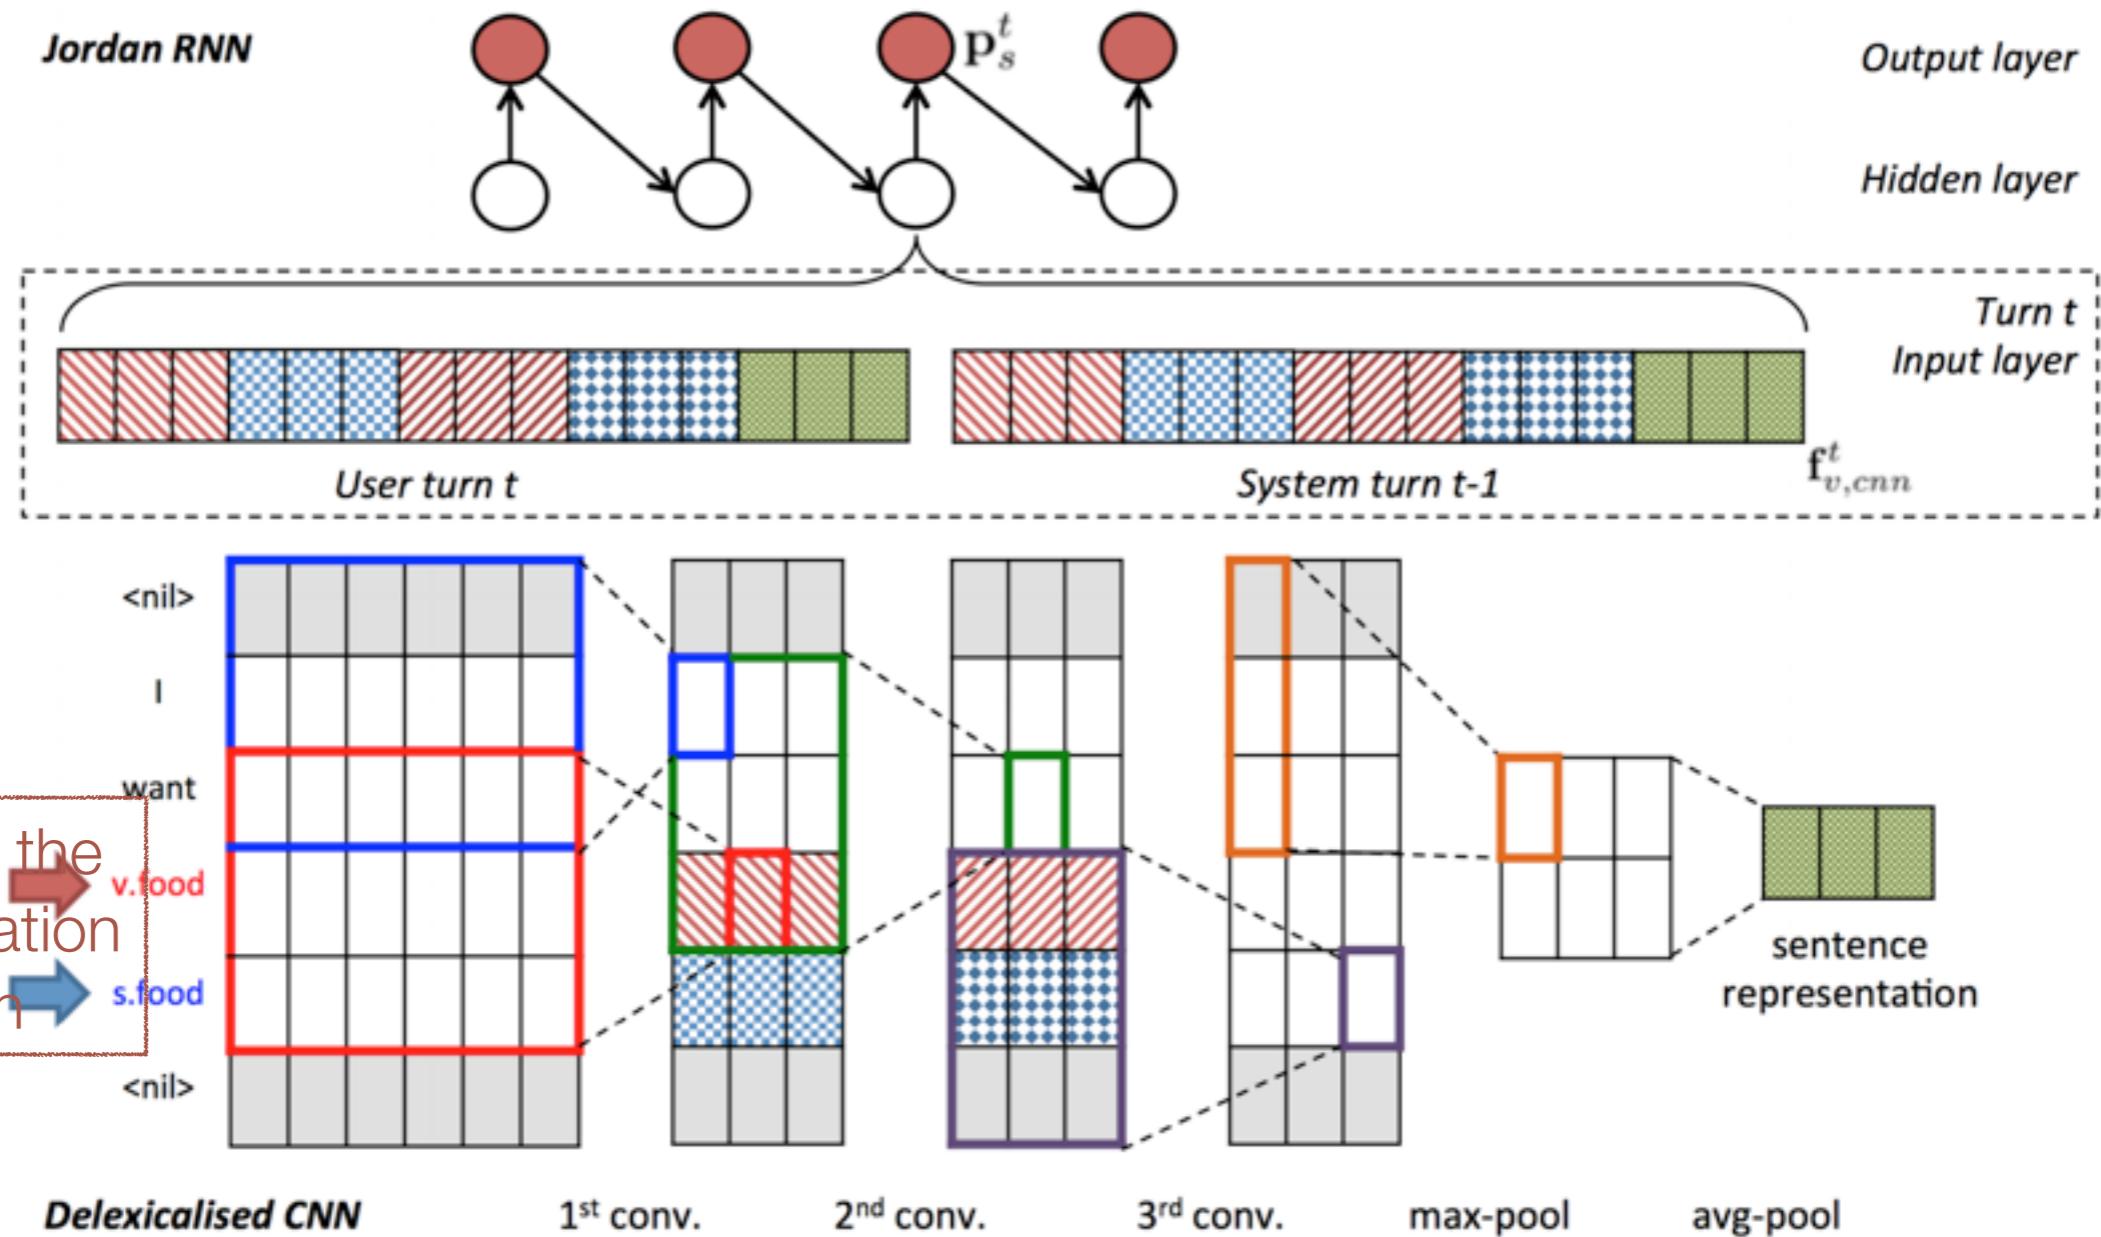
Jordan RNN--CNN Belief Trackers



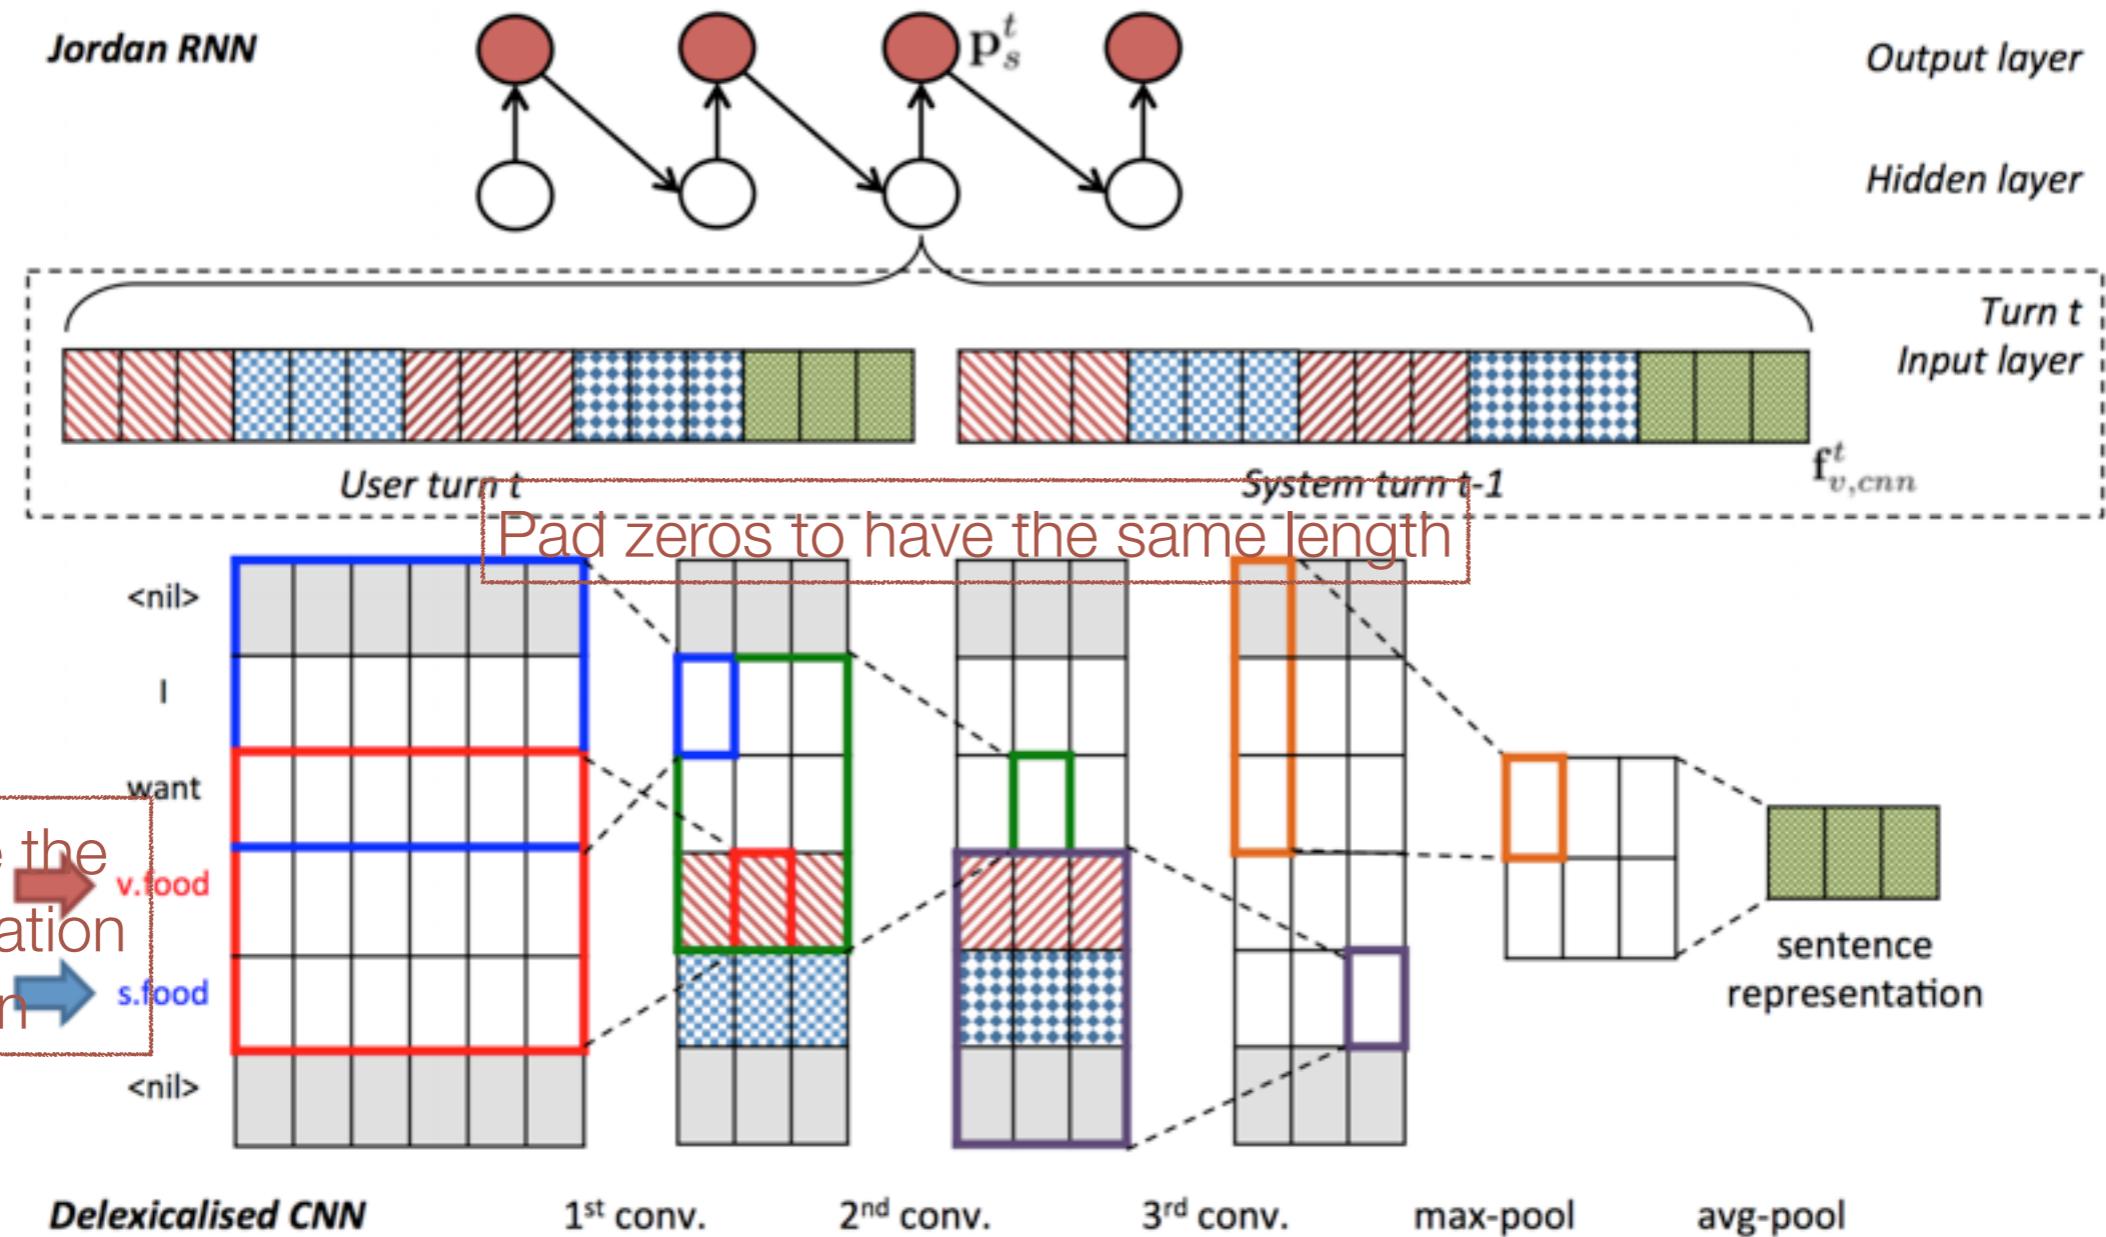
Visualize CNN [13]



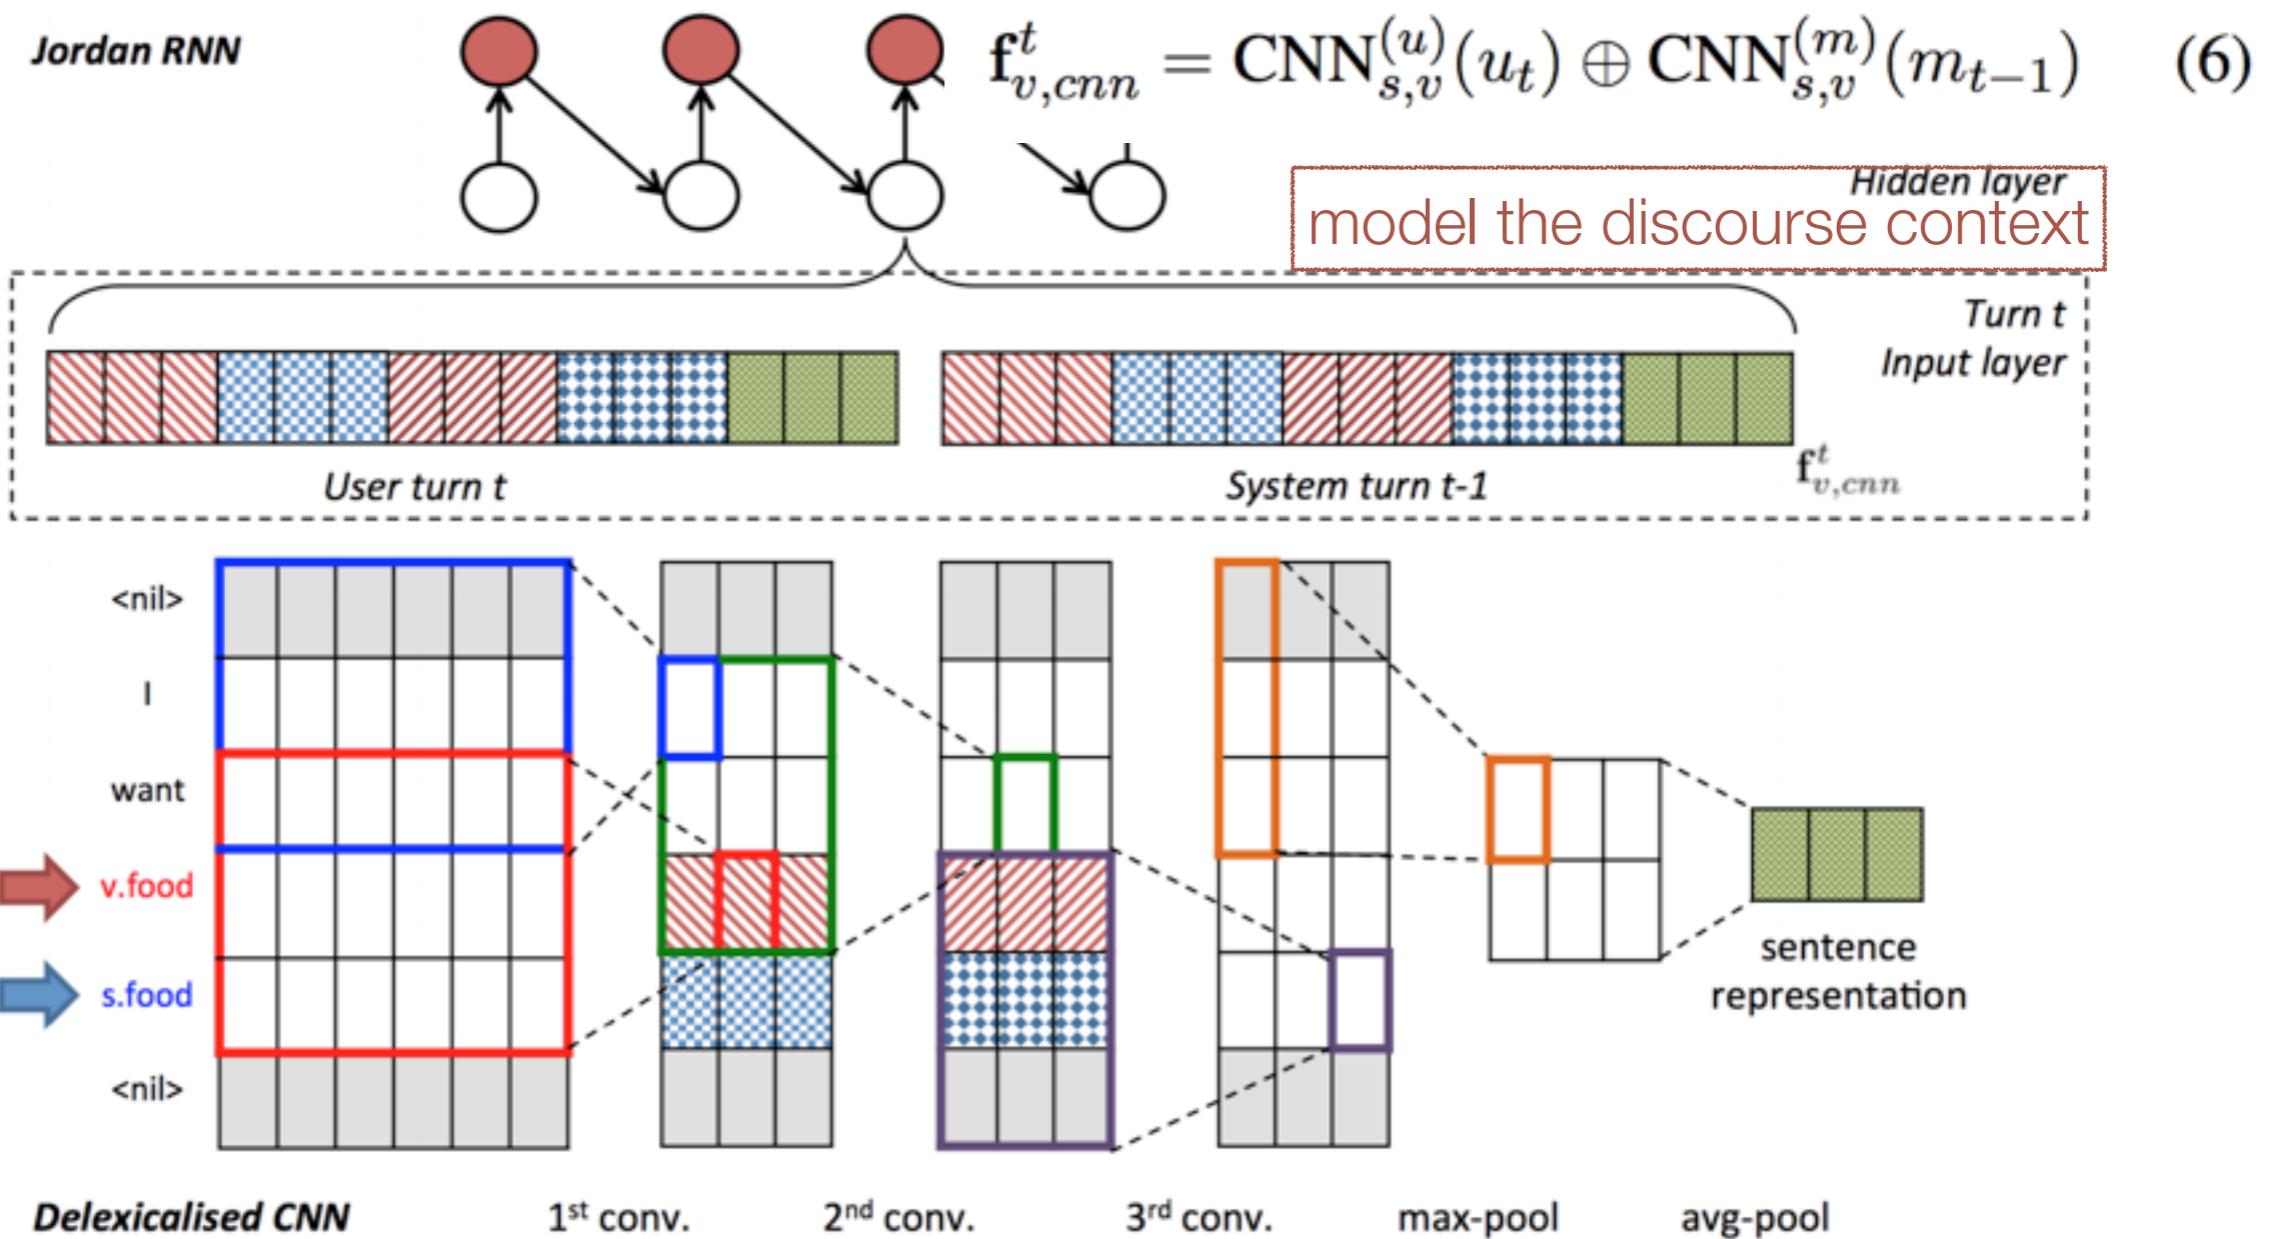
Jordan RNN--CNN Belief Trackers



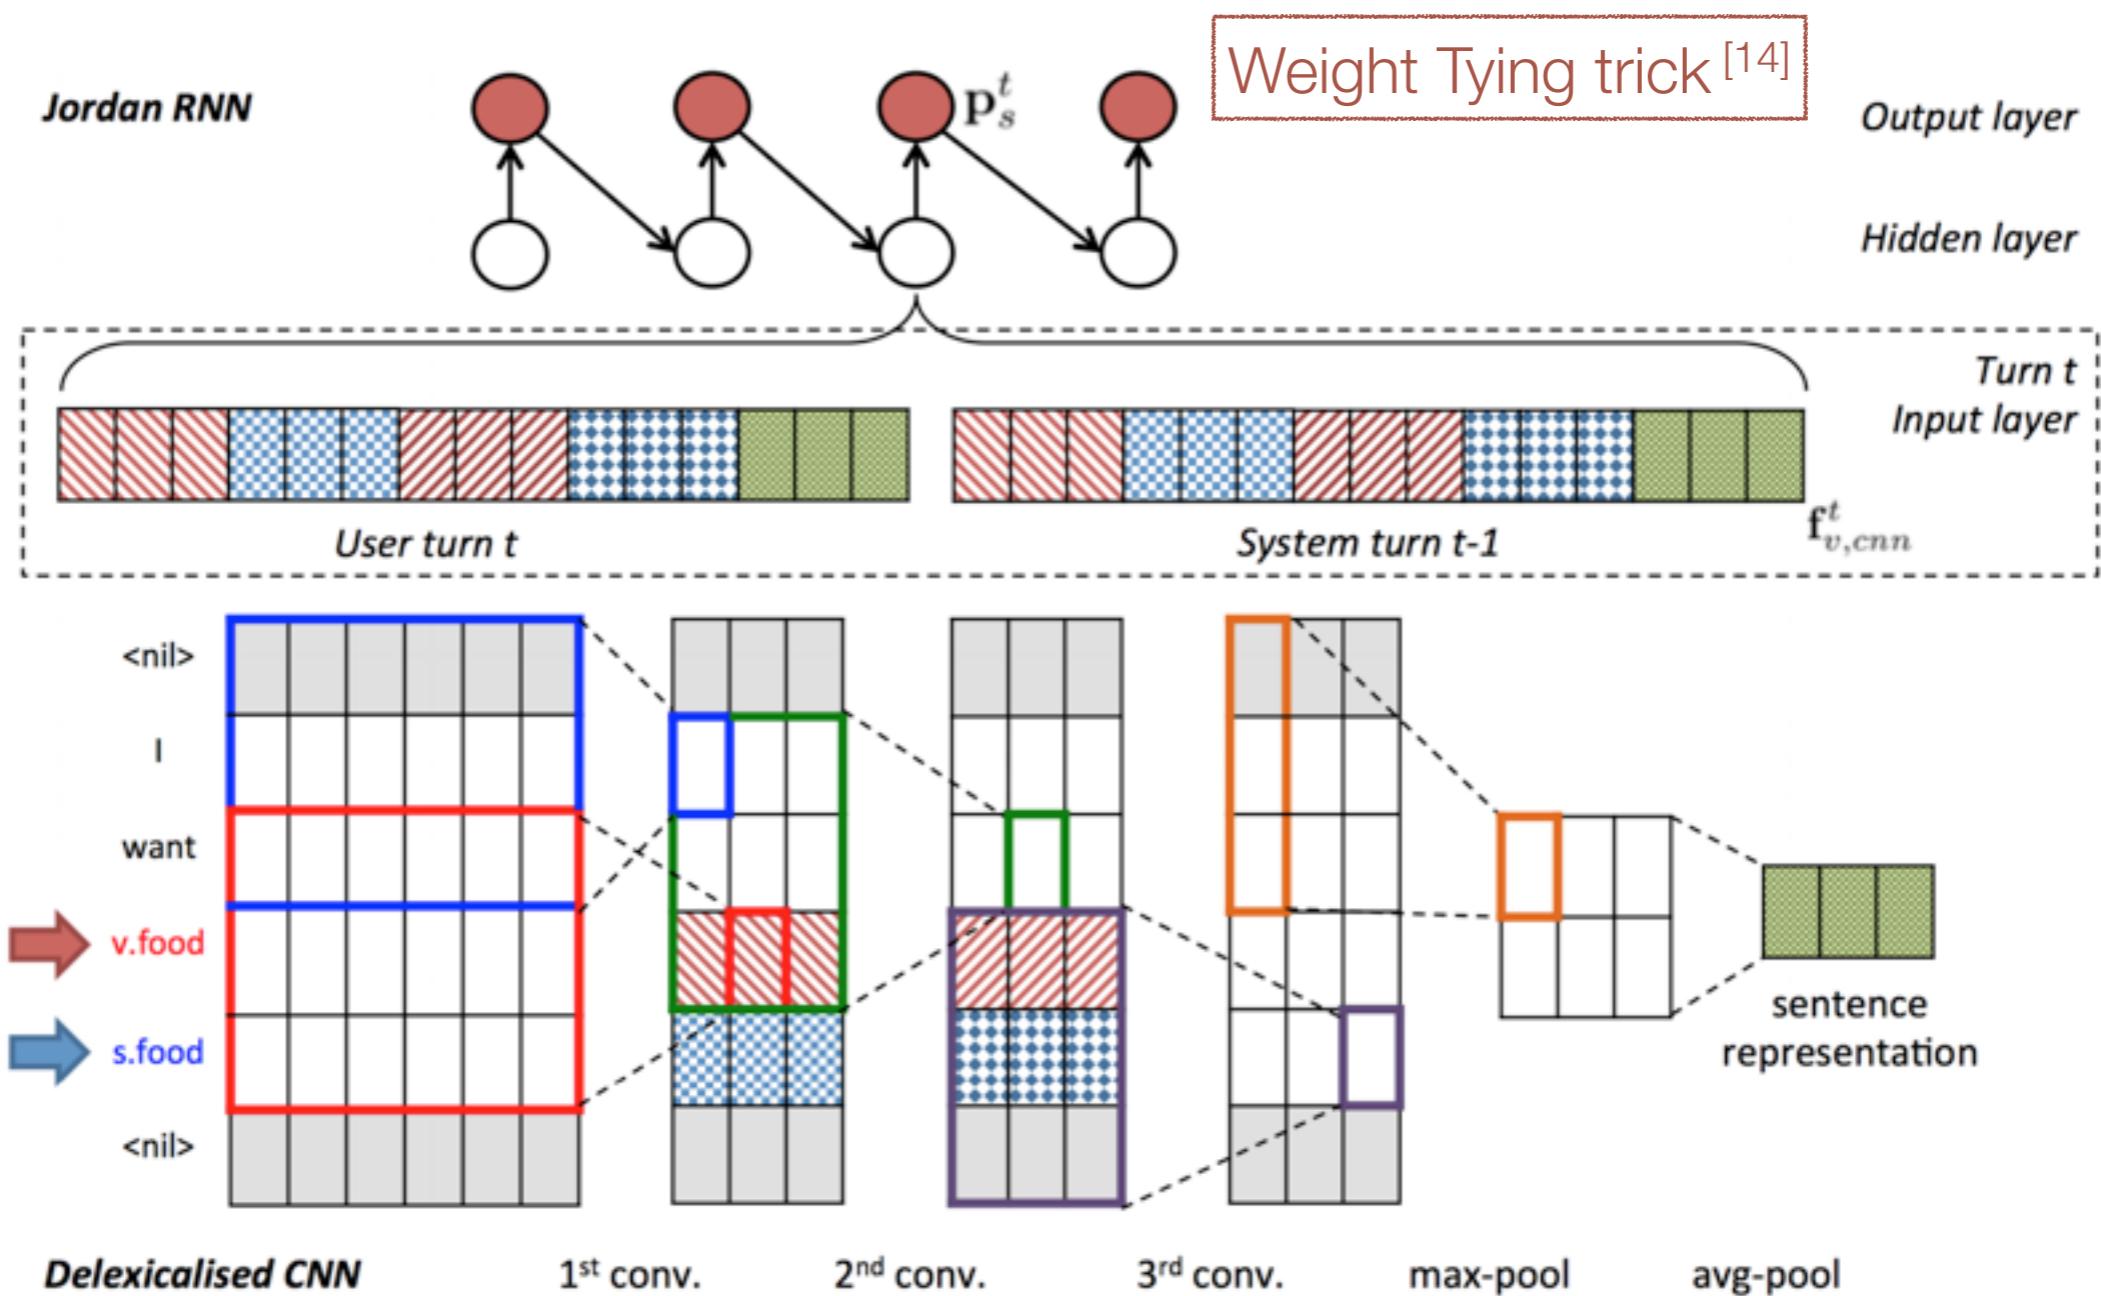
Jordan RNN--CNN Belief Trackers



Jordan RNN--CNN Belief Trackers



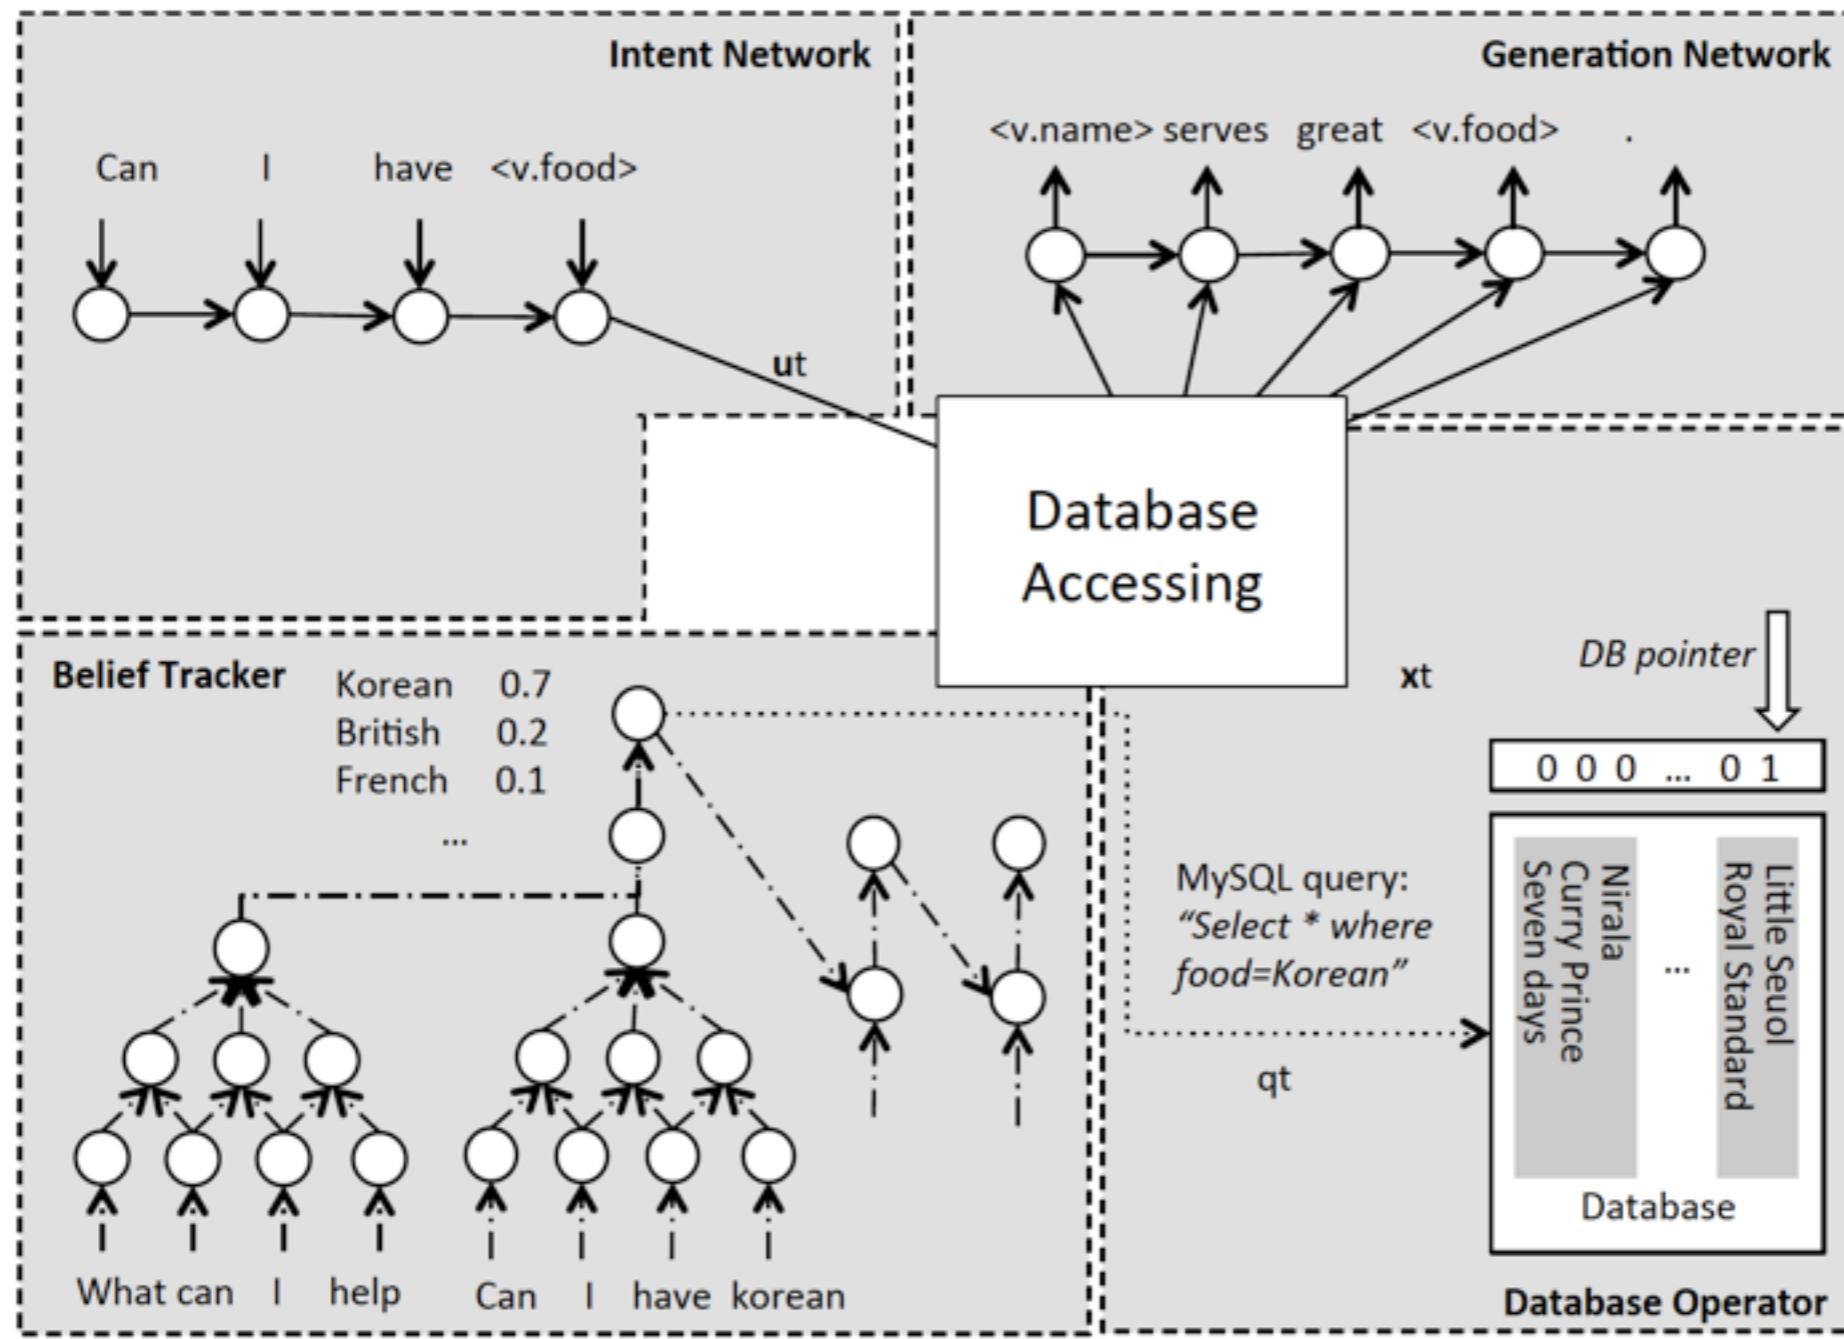
Jordan RNN--CNN Belief Trackers



Jordan RNN--CNN Belief Trackers

```
13 # Informable slot tracker
14 class CNNInformableTracker(BaseNNModule):
15
16     def __init__(self, belief_size,
17                  ivocab_size, ihidden_size,
18                  ovocab_size, ohidden_size):
19
20         # parameters for RNN tracker
21         # tracker specific CNN encoder
22         self.sCNN = CNNEncoder(ivocab_size, ihidden_size,
23                                pool=(False, False, True), level=3)
24         self.tCNN = CNNEncoder(ovocab_size, ihidden_size,
25                                pool=(False, False, True), level=3)
26
27         # handling feature to belief mapping
28         self.dbm1 = belief_size-1
29         self.db   = belief_size
30         self.dh   = int(belief_size/1.5)
31         self.Wfbs = theano.shared(0.3 * np.random.uniform(-1.0,1.0, \
32                                (ihidden_size*5, self.dh)).astype(theano.config.floatX))
33         self.Wfbt = theano.shared(0.3 * np.random.uniform(-1.0,1.0, \
34                                (ihidden_size*5, self.dh)).astype(theano.config.floatX))
35         self.Whb  = theano.shared(0.3 * np.random.uniform(-1.0,1.0, \
36                                (self.dh)).astype(theano.config.floatX))
```

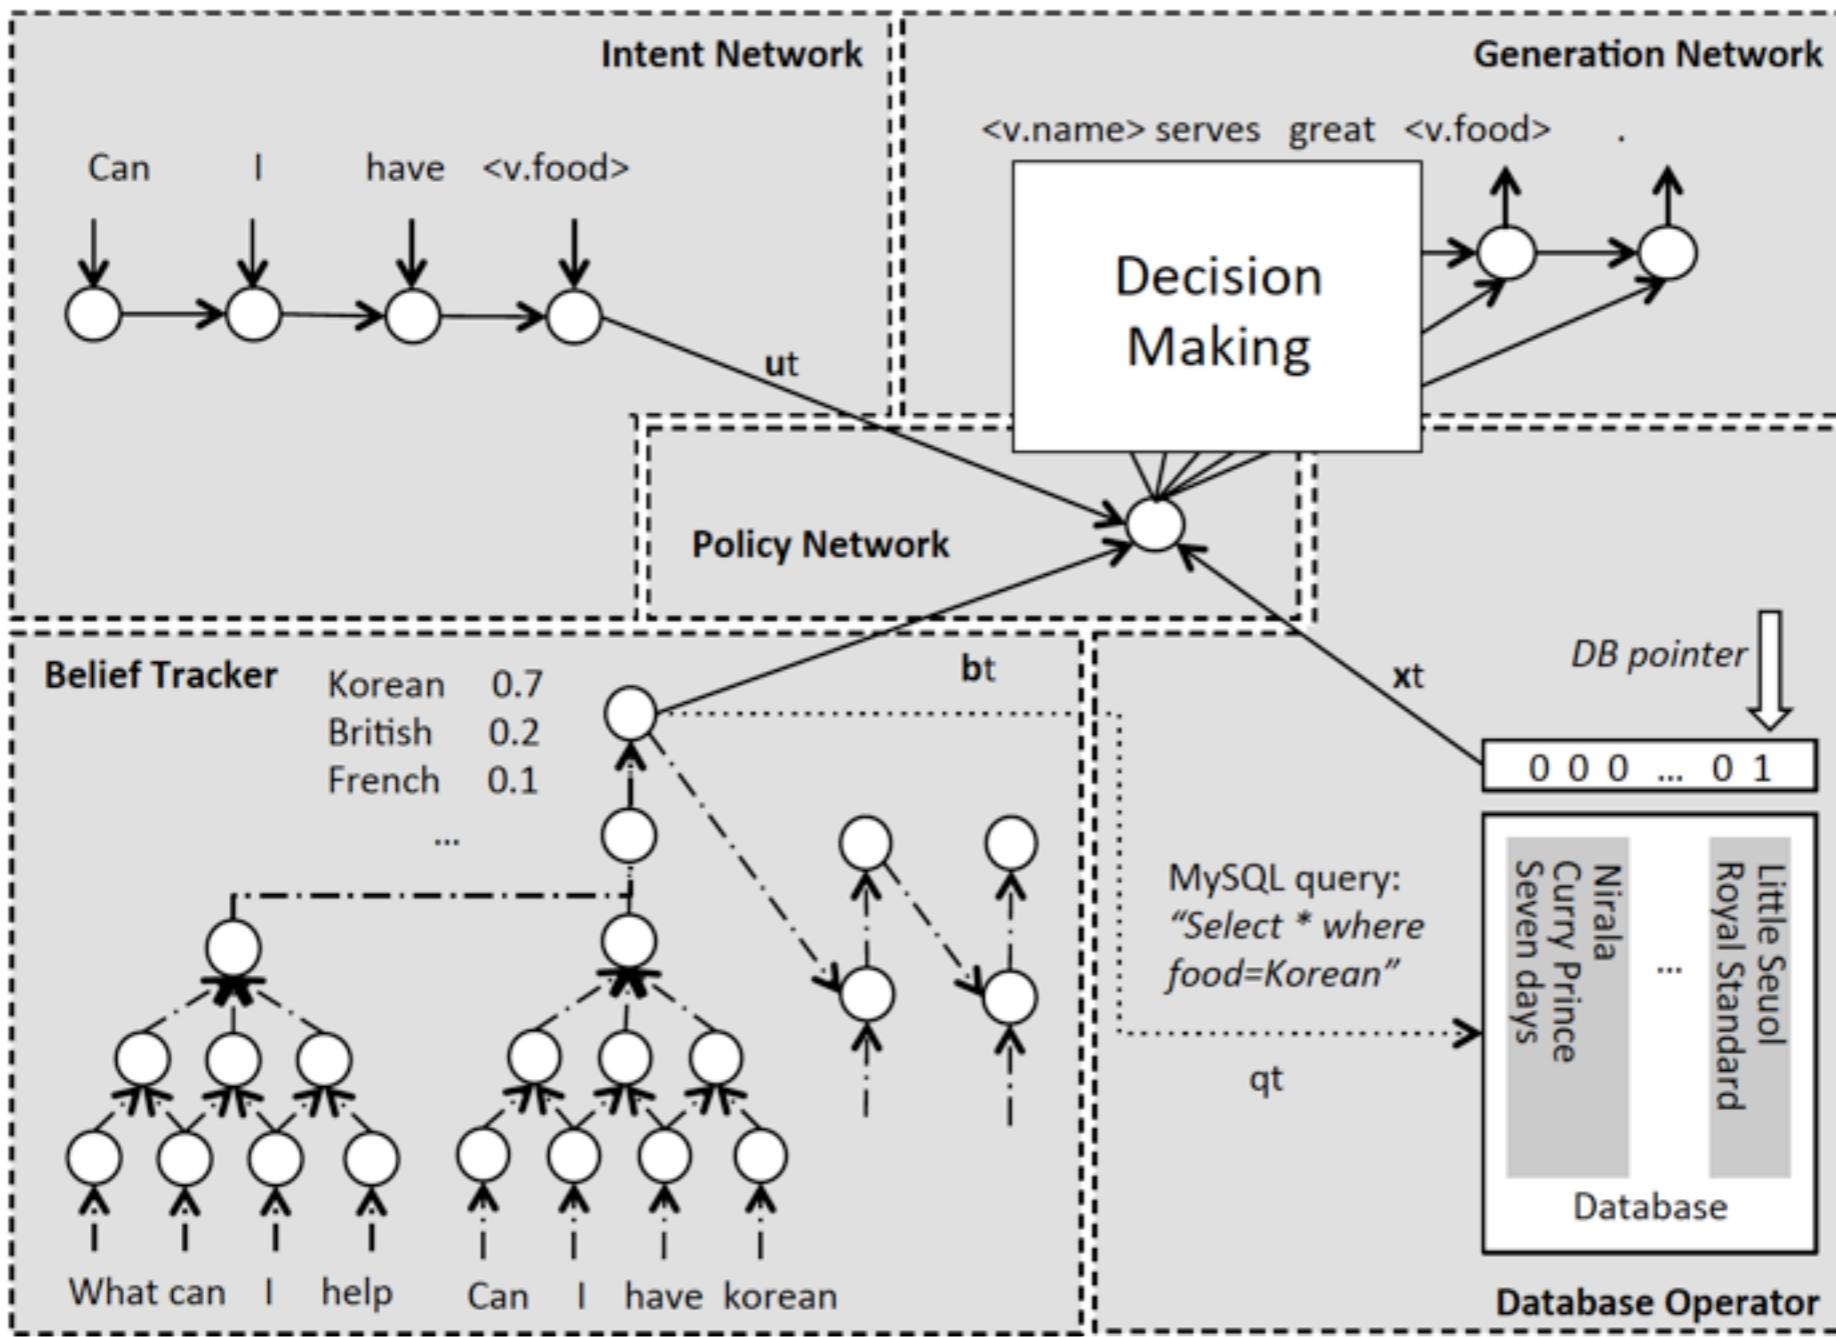
Task-oriented Neural Dialog System



Database Operator

```
948     # search database functio 289             # choose venue
949     def _searchDB(self,b):    290             venues = [i for i, e in enumerate(db_degree_t[:-6]) if e != 0 ]
950
951         query = []           291             # keep the current venue
952         q = []                292             if selected_venue in venues: pass
953         db_logic = []          293             else: # choose the first match as default index
954         # formulate query for 294                 if len(venues)!=0: selected_venue = random.choice(venues)
955         if self.trkinf==True:   295                 # no matched venues
956             for i in range(len(self.inf_dimensions)-1):
957                 b_i = b[self.inf_dimensions[i]:self.inf_dimensions[i+1]]
958                 idx = np.argmax(np.array(b_i)) + self.inf_dimensions[i]
959                 # ignore dont care case
960                 s2v = self.reader.infovs[idx]
961                 if '=dontcare' not in s2v and '=none' not in s2v:
962                     query.append(idx)
963                     q.append(idx)
964             # search through db by query
965             for entry in self.reader.db2inf:
966                 if set(entry).issuperset(set(query)):
967                     db_logic.append(1)
968                 else:
969                     db_logic.append(0)
970             # form db count features
971             dbcount = sum(db_logic)
```

Task-oriented Neural Dialog System



Policy Network

- can be viewed as a glue which outputs a single action vector
- an additional ScoreTable and additional matching degrees/counts

Scoring Table

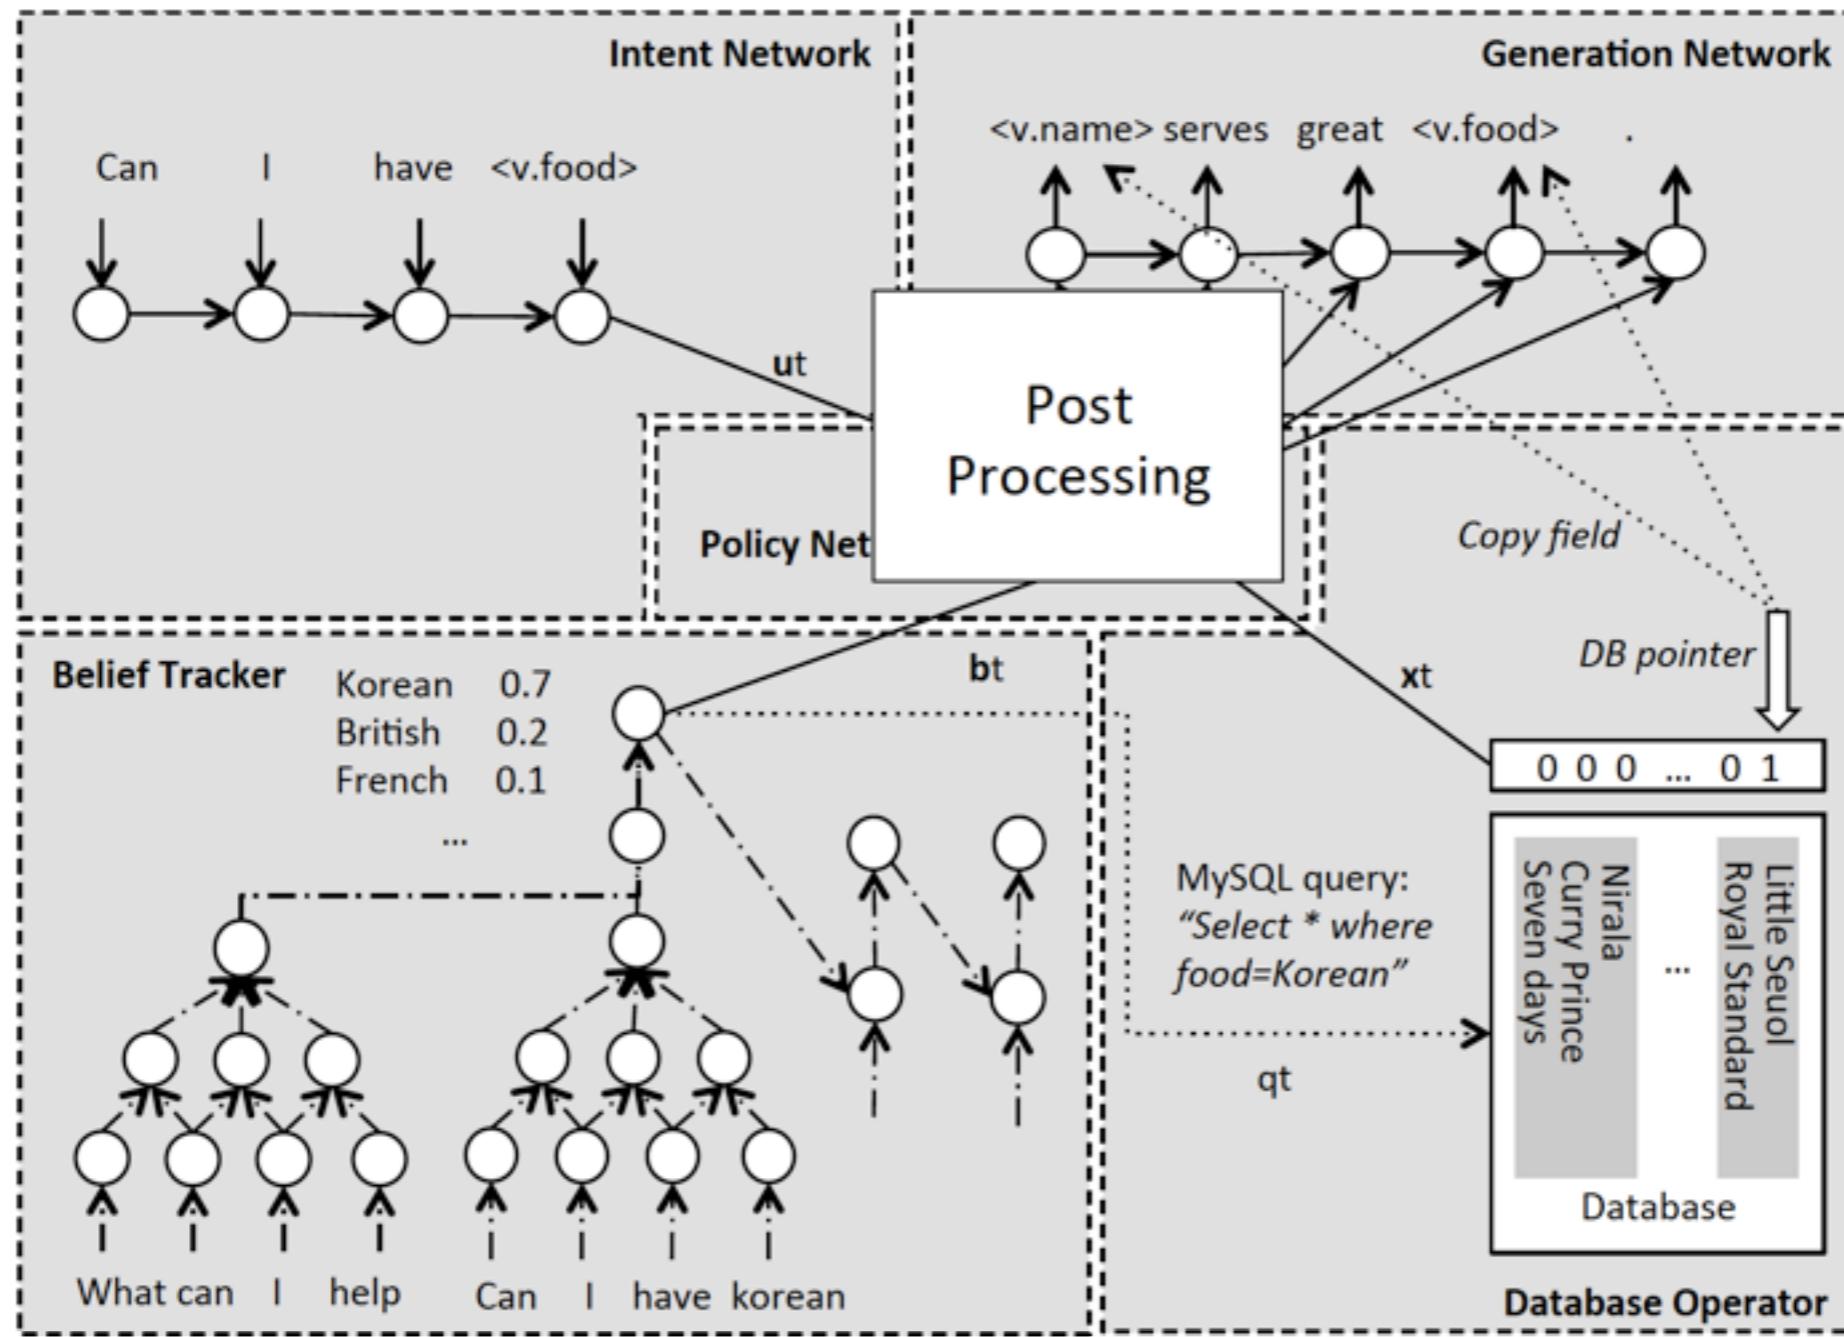
Table 5: Additional R_t term for delexicalised tokens when using weighted decoding (Equation 14). *Not observed* means the corresponding tracker has a highest probability on either *not mentioned* or *dontcare* value, while *observed* mean the highest probability is on one of the categorical values. A positive score encourages the generation of that token while a negative score discourages it.

| Delexicalised token | Examples | R_t (<i>observed</i>) | R_t (<i>not observed</i>) |
|-------------------------|----------------------------|---------------------------|-------------------------------|
| informable slot token | <s.food>, <s.area>,... | 0.0 | 0.0 |
| informable value token | <v.food>, <v.area>,... | +0.05 | -0.5 |
| requestable slot token | <s.phone>, <s.address>,... | +0.2 | 0.0 |
| requestable value token | <v.phone>, <v.address>,... | +0.2 | 0.0 |

Policy Network

```
1168     def _genScoreTable(self, sem_j) 280             # search DB
1169         scoreTable = {} 281
1170         # requestable tracker score 282             # score table
1171         if self.trk=='rnn' and self 283             scoreTable = self._genScoreTable(full_belief_t)
1172             infbn = 3 if self.trkinf 284             # generation
1173             for i in range(len(self)) 285             generated,sample_t,_ = self.model.talk(
1174                 bn = self.req_dimer 286                 masked_intent_t,belief_t, db_degree_t,
1175                 # prediction for this 287                 masked_source_t, masked_target_t, scoreTable)
1176                 psem = self.reader. 288
1177                     np.argmax(np.array(sem_j[infbn+i])) +\
1178                         self.req_dimensions[i] ]
1179
1180         #print psem
1181         # slot & value
1182         s,v = psem.split('=')
1183         if s=='name': # skip name slot
1184             continue
1185         # assign score, if exist, +reward
1186         score = -0.05 if v=='none' else 0.2
1187         # slot value indexing
1188         vidx = self.reader.vocab.index('[VALUE_'+s.upper()+']')
1189         sidx = self.reader.vocab.index('[SLOT_'+s.upper()+']')
1190         scoreTable[sidx] = score
1191         scoreTable[vidx] = score # reward [VALUE_****] if generate
1192         # informative tracker scoreTable
1193         if self.trk=='rnn' and self.trkinf==True:
```

Task-oriented Neural Dialog System



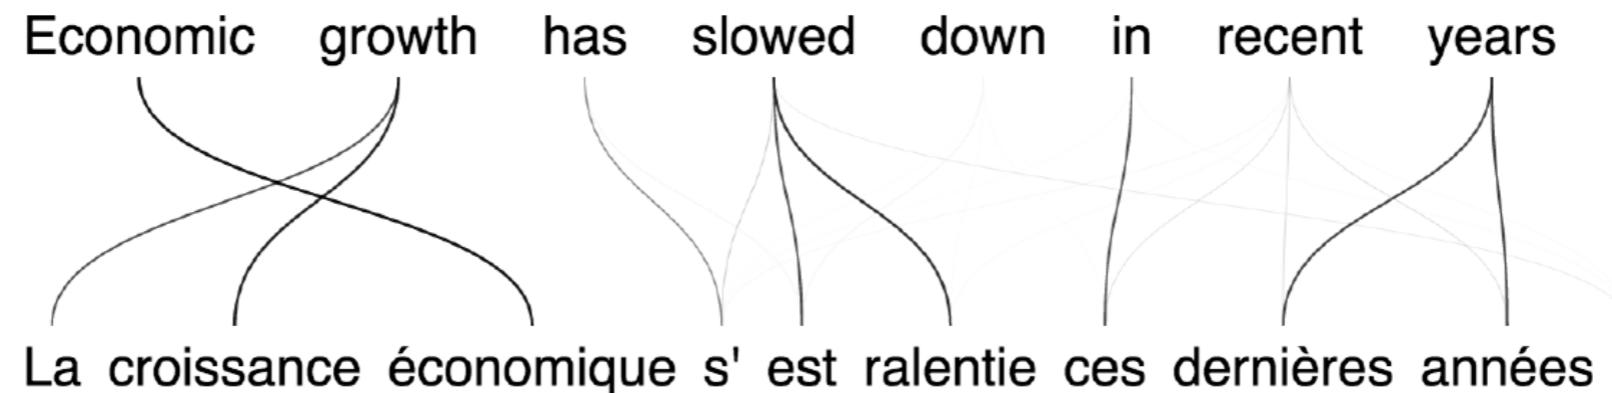
Generation Network

- generates **template-like** sentences token by token based on the language model probabilities and the system action output by policy network.
- Once the output token sequence has been generated, the generic tokens are replaced by their actual values: (1) replacing delexicalised slots by random sampling from a list of surface forms, e.g. <s.food> to food or type of food, and (2) replacing delexicalised values by the actual attribute values of the entity currently selected by the DB pointer.
[7]
- **attentive** generation network.

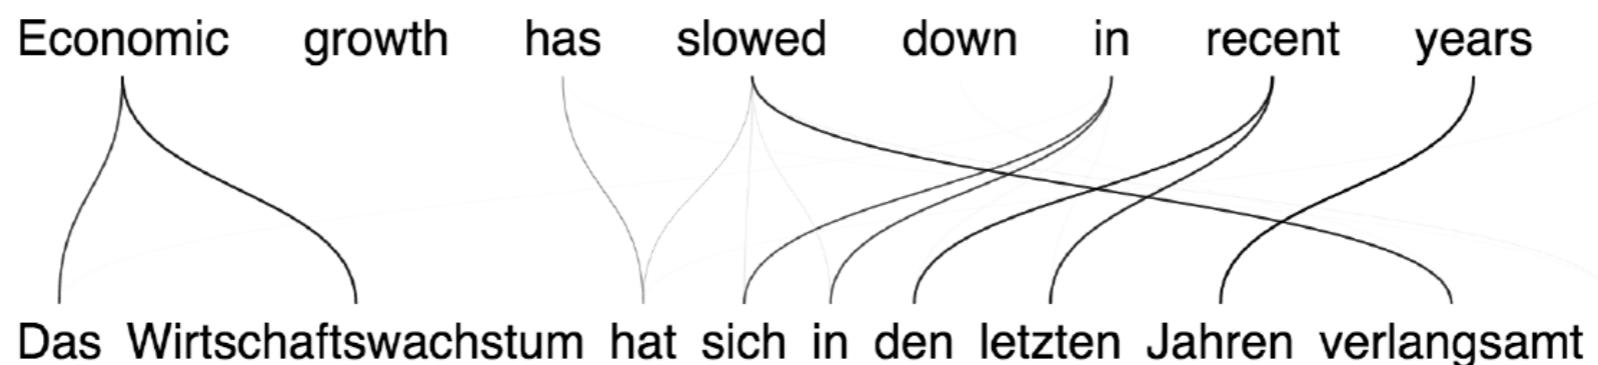
Attention-based Neural Machine Translation

- ♦ fixed representation of source sentence → soft and dynamic^[5]

English-French

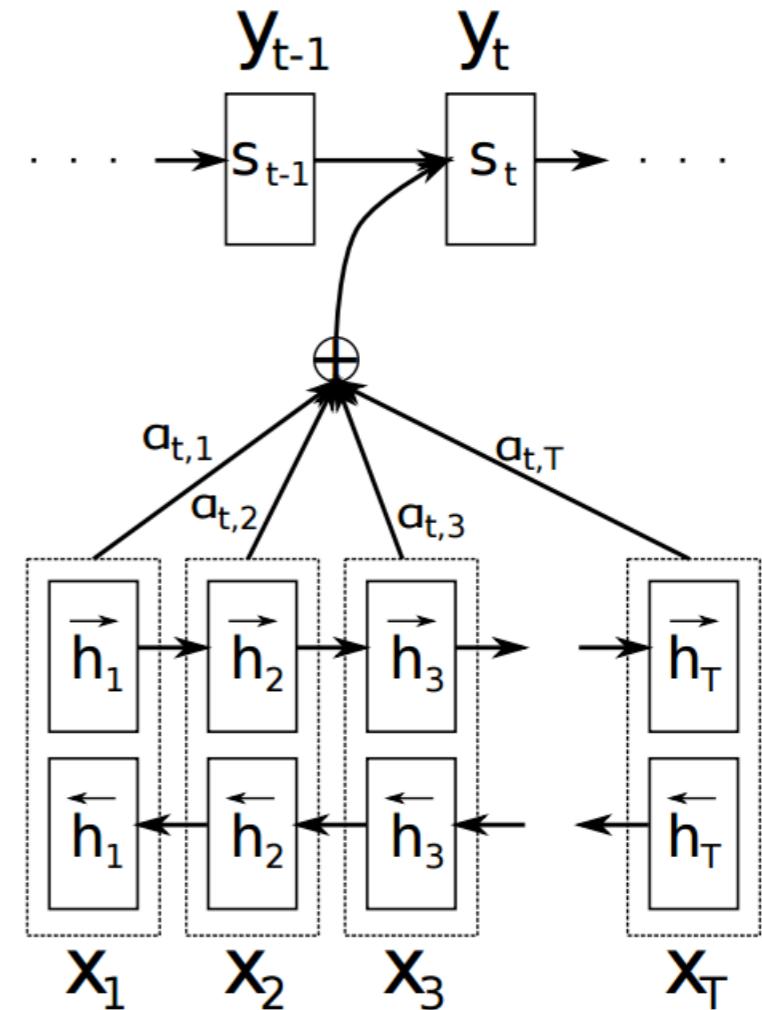


English-German



Recall Attention Mechanism

- At every generation step t [5]
 - Score source h_j by
$$e_{tj} = \mathbf{v}^T \tanh(\mathbf{W} \cdot s_{t-1} + \mathbf{U} \cdot h_j)$$
$$\alpha_{tj} = \text{softmax}(e_{tj})$$
 - Take an expectation over sources
$$c_t = \sum_j \alpha_{tj} h_j$$
 - Everything is differentiable.
Back-prop end-to-end!



Attentive Generation Network

- use attention mechanism to combine the tracker belief states. In fact, attentive policy.

$$e_{tj} = \mathbf{v}^T \tanh(\mathbf{W} \cdot \mathbf{s}_{t-1} + \mathbf{U} \cdot \mathbf{h}_j)$$

$$\alpha_{tj} = \text{softmax}(e_{tj})$$

$$\mathbf{o}_t^{(j)} = \tanh(\mathbf{W}_{zo} \mathbf{z}_t + \hat{\mathbf{p}}_t^{(j)} + \mathbf{W}_{xo} \hat{\mathbf{x}}_t) \quad (10)$$

where for a given ontology \mathbb{G} ,

$$\hat{\mathbf{p}}_t^{(j)} = \sum_{s \in \mathbb{G}} \alpha_s^{(j)} \tanh(\mathbf{W}_{po}^s \cdot \hat{\mathbf{p}}_s^t) \quad (11)$$

and where the attention weights $\alpha_s^{(j)}$ are calculated by a scoring function,

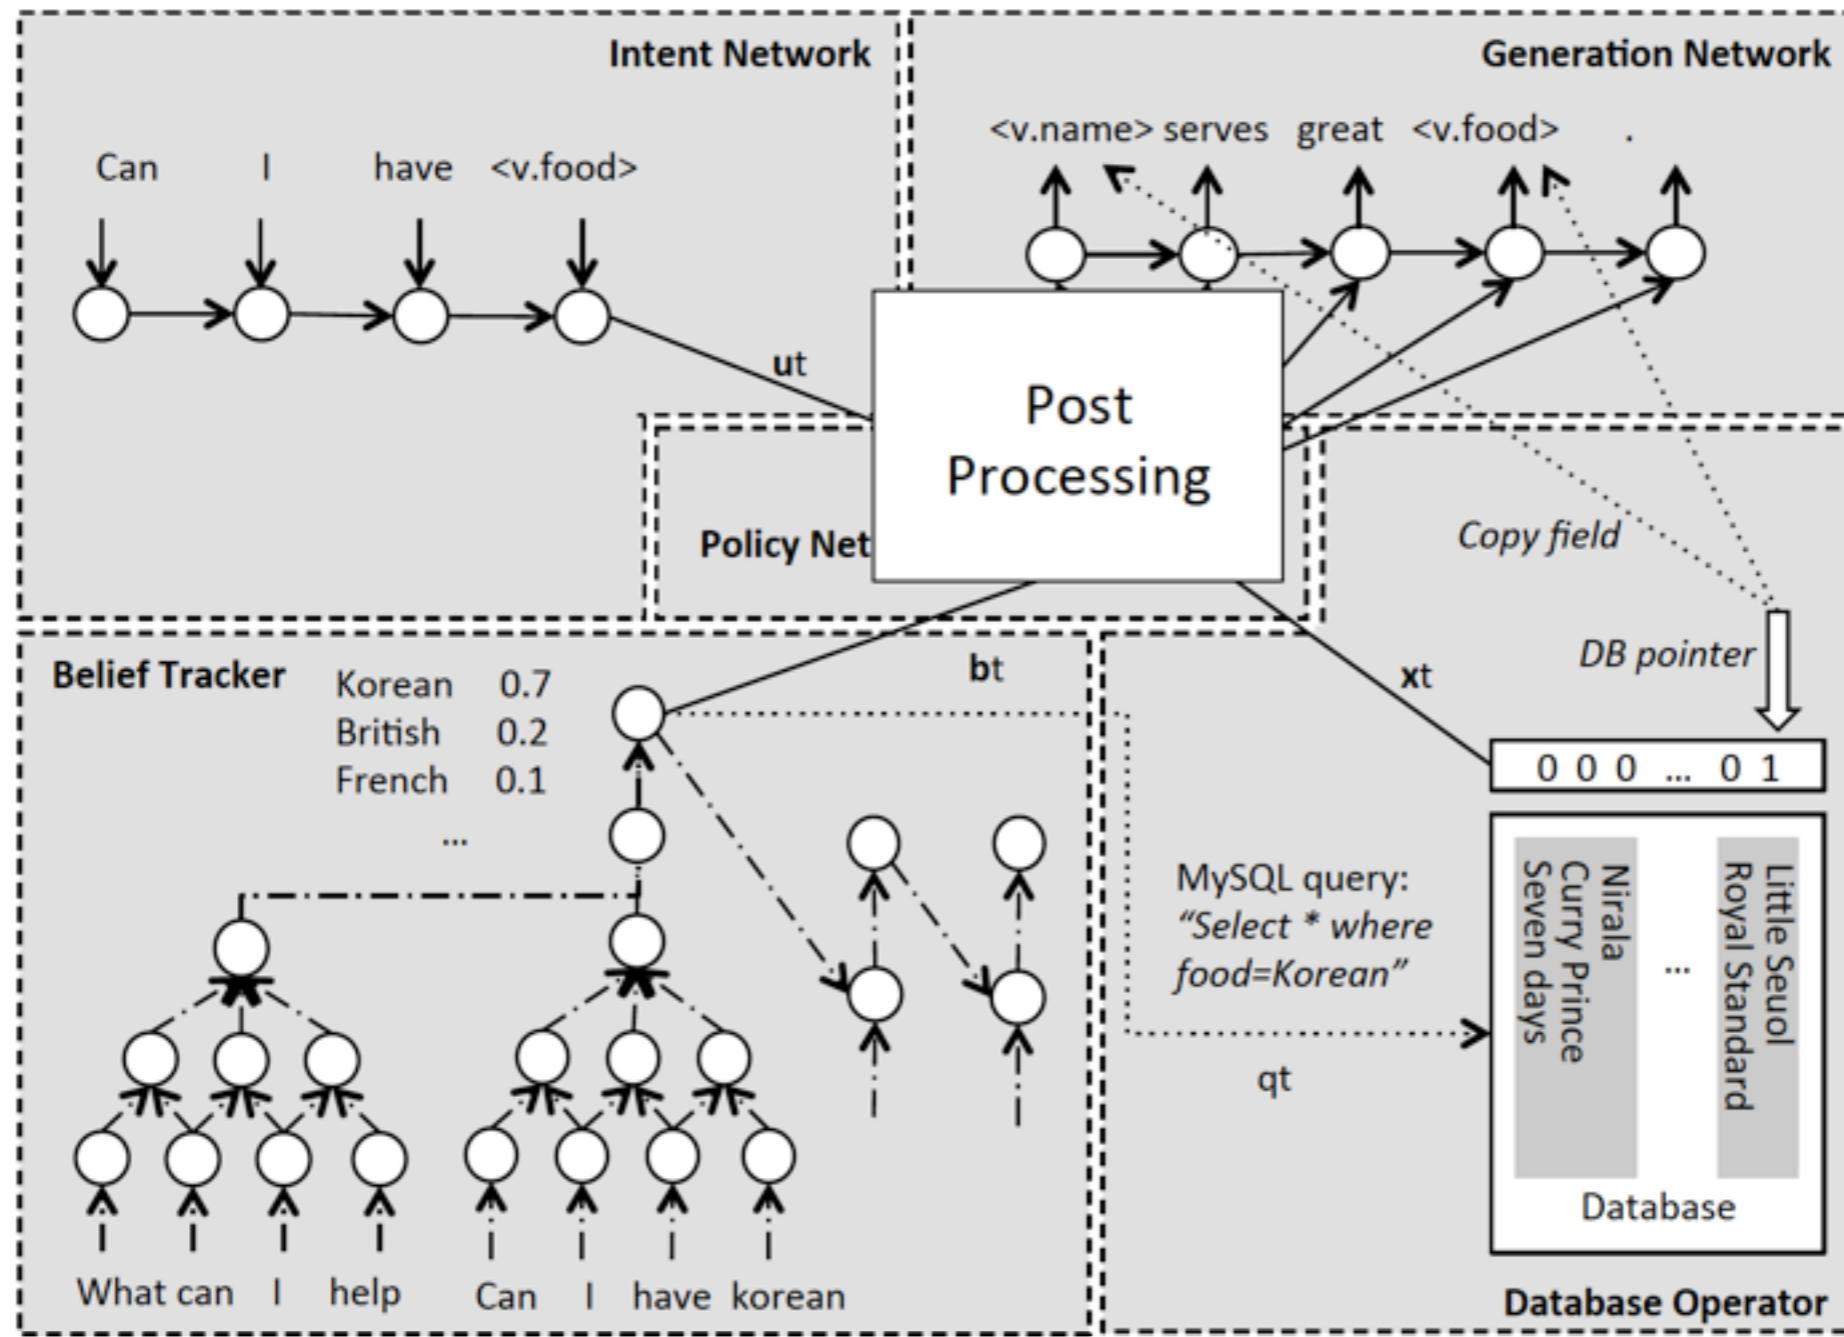
$$\alpha_s^{(j)} = \text{softmax}(\mathbf{r}^T \tanh(\mathbf{W}_r \cdot \mathbf{u}_t)) \quad (12)$$

where $\mathbf{u}_t = \mathbf{z}_t \oplus \hat{\mathbf{x}}_t \oplus \hat{\mathbf{p}}_s^t \oplus \mathbf{w}_j^t \oplus \mathbf{h}_{j-1}^t$, matrix \mathbf{W}_r , and vector \mathbf{r} are parameters to learn and \mathbf{w}_j^t is the embedding of token w_j^t .

Attentive Generation Network

```
402     def prepareBelief(self,belief_t):
403         # belief vectors
404         beliefs_t = []
405         bn = 0
406         for bvec in belief_t:
407             size = bvec.shape[0]
408             beliefs_t.append( T.tanh(T.dot(bvec,self.Ws1[bn:bn+size,:])).\
409                               dimshuffle('x',0) )
410             bn += size
411         beliefs_t = T.concatenate(beliefs_t, axis=0)
412         return beliefs_t
413
414
415     # decoding function
416     def decode(self, masked_source_t, masked_source_len_t,
417               masked_target_t, masked_target_len_t,
418               intent_t, belief_t, degree_t, utt_group_t,
419               snapshot_t, sample_t):
420         # decide policy type
421         if self.ply=='attention':
422             # set dummpy, not used
423             actEmb_t = T.concatenate(belief_t, axis=0)
424             prior_t = posterior_t = z_t = base_t = debug_t = None
425             # attentive policy
426             belief_t = self.policy.prepareBelief(belief_t)
```

Task-oriented Neural Dialog System



Task-oriented Neural Dialog System

Task-oriented Neural Dialog System

```
274     # read and understand user sentence
275     masked_intent_t = self.model.read(masked_source_t)
276     full_belief_t, belief_t = self.model.track(
277         flatten_belief_tm1, masked_source_t, masked_target_tm1,
278         srcfeat_t, tarfeat_tm1 )
279     flatten_belief_t = np.concatenate(full_belief_t, axis=0)
280     # search DB
281     db_degree_t, query = self._searchDB(flatten_belief_t)
282     # score table
283     scoreTable = self._genScoreTable(full_belief_t)
284     # generation
285     generated, sample_t, _ = self.model.talk(
286         masked_intent_t, belief_t, db_degree_t,
287         masked_source_t, masked_target_t, scoreTable)
288
289     # choose venue
290     venues = [i for i, e in enumerate(db_degree_t[:-6]) if e != 0 ]
291     # keep the current venue
292     if selected_venue in venues: pass
293     else: # choose the first match as default index
294         if len(venues)!=0: selected_venue = random.choice(venues)
295         # no matched venues
296         else: selected_venue = None
```

Intent Network

Belief Tracker

DB Pointer

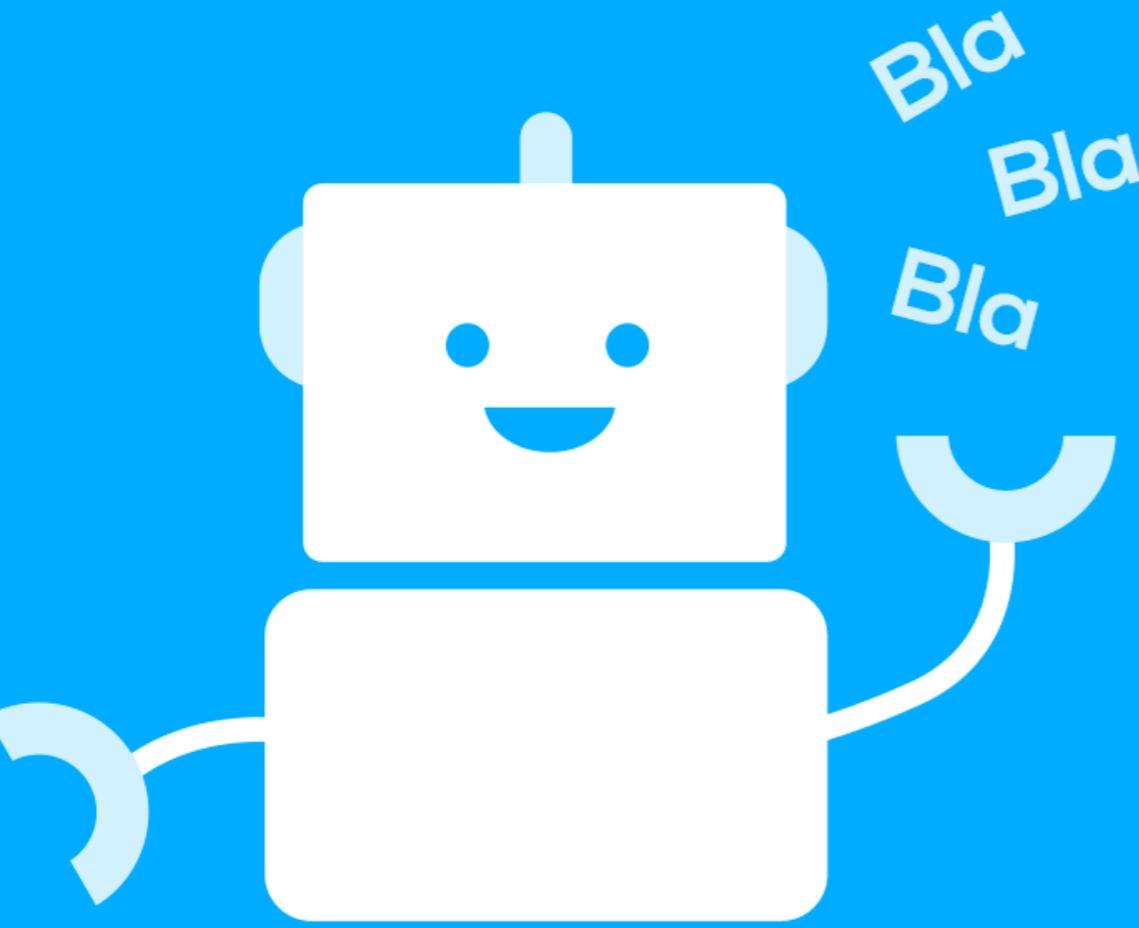
Policy Network and
Generation Network

```

298     # lexicalise generated utterance
299     generated_utts = []
300     for gen in generated:
301         generated_utt = ' '.join([self.reader.vocab[g] for g in gen[0]])
302         generated_utts.append(generated_utt)
303     gennerated_utt = generated_utts[0]
304
305     # calculate semantic match rate
306     twords = [self.reader.vocab[w] for w in masked_target_t]
307     for gen in generated:
308         gwords = [self.reader.vocab[g] for g in gen[0]]
309         for gw in gwords:
310             if gw.startswith('[VALUE_') or gw.startswith('[SLOT_'):
311                 if gw in twords: # match target semi token
312                     stats['approp'][0] += 1.0
313                     stats['approp'][1] += 1.0
314             #gstats += np.mean(np.array(gen[2][1:]),axis=0 )
315             num_sent += 1
316
317     # update history belief
318     flatten_belief_tm1 = flatten_belief_t[:self.inf_dimensions[-1]]
319
320     # for calculating success: check requestable slots match
321     requestables = ['phone','address','postcode','food','area','pricerange']
322     for requestable in requestables:
323         if '[VALUE_'+requestable.upper()+']' in gennerated_utt:
324             reqs.append(self.reader.reqs.index(requestable+'=exist'))
325     # check offered venue
326     if '[VALUE_NAME]' in generated_utt and selected_venue!=None:
327         venue_offered = self.reader.db2inf[selected_venue]

```

Policy Network and Generation Network



Bla
Bla
Bla

vs.



Task-oriented

Dataset

CamRest676 Dataset

- <https://www.repository.cam.ac.uk/handle/1810/260970>
- Cambridge restaurant dialogue domain. In-domain human-human conversation.
- Users can look for restaurant with particular food type, price range, and area (informable), and ask for phone number, address and postcode (requestable).
- Including coarse slot-value labels per turn.
- Evaluation: *Dialogue success, BLEU

CamRest676 Dataset

Task:
Find a **restaurant**,
cheap, Korean, North
Ask **phone number**



*Hi, I want a **cheap Korean restaurant**.*

*What **area** are you looking for ?*

*Somewhere in the **north**.*

Little Seoul is nice one in the north.

*What is the **phone number**?*

*Its phone number is **01223456789**.*

Thank you very much, good bye.

Thank you for using the system.



What user wants?

| | |
|-------|--------|
| Food | Korean |
| Price | Cheap |
| Area | North |

Search Table

| | |
|--------------|-----|
| Little Seoul | ... |
|--------------|-----|

Dataset Statistics

- Ontology:
 - Cambridge restaurant domain, 99 venues.
 - 3 informative slots: area, price range, food type
 - 3 requestable slots: address, phone, postcode
- Dataset:
 - 676 dialogues, ~2750 turns
 - 3000 HITS, takes 3 days, costs ~400 USD
 - Data cleaning takes 2-3 days for one person

Training Details

- Two stage training:
 - Training trackers with label cross entropy
 - Training other parts with response cross entropy
- Decoding
 - Beam search w/ beam width 10
 - Decode with average word likelihood

Experiment Results

Table 1: Tracker performance in terms of Precision, Recall, and F-1 score.

| Tracker type | Informable | | | Requestable | | |
|-----------------|------------|--------|--------|-------------|--------|--------|
| | Prec. | Recall | F-1 | Prec. | Recall | F-1 |
| cnn | 99.77% | 96.09% | 97.89% | 98.66% | 93.79% | 96.16% |
| ngram | 99.34% | 94.42% | 96.82% | 98.56% | 90.14% | 94.16% |

Table 3: Human assessment of the NN system. The rating for comprehension/naturalness are both out of 5.

| Metric | NN |
|-----------------|------|
| Success | 98% |
| Comprehension | 4.11 |
| Naturalness | 4.05 |
| # of dialogues: | 245 |

Example Dialogues

Human : hello
Machine : what cuisine would you like ?
Human : afghan
Machine : i am sorry , there are no afghan restaurants . would you like something else ?
Human : korean
Machine : little seoul serves korean food , it is located at 108 regent street city centre .
Human : is it expensive?
Machine : little seoul serves korean food and is in the expensive price range .
Human : what area is that
Machine : little seoul is located in the centre area of town . would you like their number ?
Human : yes
Machine : the telephone number is 01223 308681 . is there anything else i can help you with today ?
Human : no goodbye
Machine : thank you for using the cambridge restaurant system , goodbye .

Homework

- ◆ Paper Reading:
 - ◆ Lifeng Shang, Zhengdong Lu, Hang Li. “Neural Responding Machine for Short-Text Conversation”. EMNLP 2015.
 - ◆ Zongcheng Ji, Zhengdong Lu, Hang Li. “An Information Retrieval Approach to Short Text Conversation”. arXiv preprint 2014.
- ◆ Code Reading:
 - ◆ Multi-turn Response Selection: <https://github.com/MarkWuNLP/MultiTurnResponseSelection>
 - ◆ A Neural Conversational Model: <https://github.com/Conchylicultor/DeepQA>
 - ◆ Attention-based Conversational Model: https://github.com/pbhatia243/Neural_Conversation_Models



Thanks for your attention!

Q&A

References

- [1] Pascanu, Razvan, Tomas Mikolov, and Yoshua Bengio. "On the difficulty of training recurrent neural networks." International Conference on Machine Learning. 2013.
- [2] Mikolov, Tomas, et al. "Recurrent neural network based language model." Interspeech. Vol. 2. 2010.
- [3] Colah. "Understanding LSTMs". <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- [4] Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le. "Sequence to sequence learning with neural networks." Advances in neural information processing systems. 2014.
- [5] Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. "Neural machine translation by jointly learning to align and translate." arXiv preprint arXiv:1409.0473 (2014).
- [6] Vinyals, Oriol, and Quoc Le. "A neural conversational model." arXiv preprint arXiv:1506.05869 (2015).
- [7] Wen, Tsung-Hsien, et al. "A network-based end-to-end trainable task-oriented dialogue system." EACL 2017.
- [8] Bowman, Samuel R., et al. "Generating sentences from a continuous space." arXiv preprint arXiv:1511.06349 (2015).

References

- [9] Hu, Baotian, et al. "Convolutional neural network architectures for matching natural language sentences." Advances in neural information processing systems. 2014.
- [10] Qiu, Xipeng, and Xuanjing Huang. "Convolutional Neural Tensor Network Architecture for Community-Based Question Answering." IJCAI. 2015.
- [11] Wu, Yu, et al. "Sequential Match Network: A New Architecture for Multi-turn Response Selection in Retrieval-based Chatbots." arXiv preprint arXiv:1612.01627 (2016).
- [12] Matthew Henderson, Blaise Thomson, and Steve Young. 2014. Robust dialog state tracking using delexicalised recurrent neural networks and unsupervised adaptation. In Proceedings of IEEE Spoken Language Technology.
- [13] <http://cs231n.github.io/convolutional-networks/>
- [14] Nikola Mrkšić, Diarmuid Ó Séaghdha, Blaise Thomson, Milica Gašić, Pei-Hao Su, David Vandyke, Tsung- Hsien Wen, and Steve Young. 2015. Multi-domain dialog state tracking using recurrent neural networks. In ACL, pages 794–799, Beijing, China, July. ACL.